

Lecture 2

Agents and Problem Solving

Lusi Li

**Department of Computer Science
ODU**


Reading for This Class:
Chapter 3, Russell and Norvig

Review

- **Last Class**
 - agents perceive and act in an environment
 - Description of the task environment (PEAS)
 - ideal agents maximize their performance measure
 - Rationality
 - basic agent types
 - simple reflex
 - model-based reflex
 - goal-based
 - utility-based
 - learning
- **This Class**
 - Problem Solving Agent
 - Problem Formulation
- **Next Class**
 - Problem Solving Algorithms

Review: Intelligent Agent

- **Agent: Definition**
 - Any entity that perceives its environment through sensors and acts upon that environment through actuators
 - Examples: human, robotic, software agents
- **Perception**
 - Signal from environment
- **Sensors**
 - Acquire percepts
 - Possible limitations
- **Action**
 - Attempts to affect environment
- **Actuators**
 - Transmits actions
 - Possible limitations



Agents

- **Rational Agents**
- **Agent PEAS Description**
- **Agent Environment**
 - fully observable vs. partially observable
 - deterministic vs. stochastic (non-deterministic)
 - episodic vs. sequential (non-episodic)
 - static vs. dynamic
 - discrete vs. continuous
 - single vs. multiple agents
- **Agent Types**
 - simple reflex
 - model-based reflex
 - goal-based
 - utility-based
 - learning

Problem-Solving Agents: Definition

- **Problem-solving Agent**
 - A kind of goal-based agent
 - It solves a problem by finding a sequence of actions that lead to desirable states (goals)
- **Problem solving process has four steps**
 - Formulate a “goal”
 - Formulate a “problem”
 - Search for a “solution”
 - Execute the solution

Problem-solving Agent

Preliminary Design

- **Goal Formulation**
 - Goals organize the agent behavior by limiting the objectives and hence the actions to be considered
 - A goal is formulated as **a set of world states** based on the current situation and performance measure, where the goal test is satisfied
 - Here we consider environments that are observable, discrete and deterministic
- **Problem Formulation**
 - To decide what **actions and states** to consider, given a goal
 - States: various world states (initial, goal, intermediate, etc.)
 - Actions: causing transitions between world states
- **Search**
 - There are many paths to achieve the same goal
 - The process of finding a sequence of actions that reaches the goal
- **Execution**
 - Execute the actions in the solution, one at a time

A Simple Problem-Solving Agent

function SIMPLE-PROBLEM-SOLVING-AGENT(*percept*) **returns** an action

persistent: *seq*, an action sequence, initially empty

state, some description of the current world state

goal, a goal, initially null

problem, a problem formulation

state \leftarrow UPDATE-STATE(*state*, *percept*)

if *seq* is empty **then**

goal \leftarrow FORMULATE-GOAL(*state*)

problem \leftarrow FORMULATE-PROBLEM(*state*, *goal*)

seq \leftarrow SEARCH(*problem*)

if *seq* = *failure* **then return** a null action

action \leftarrow FIRST(*seq*)

seq \leftarrow REST(*seq*)

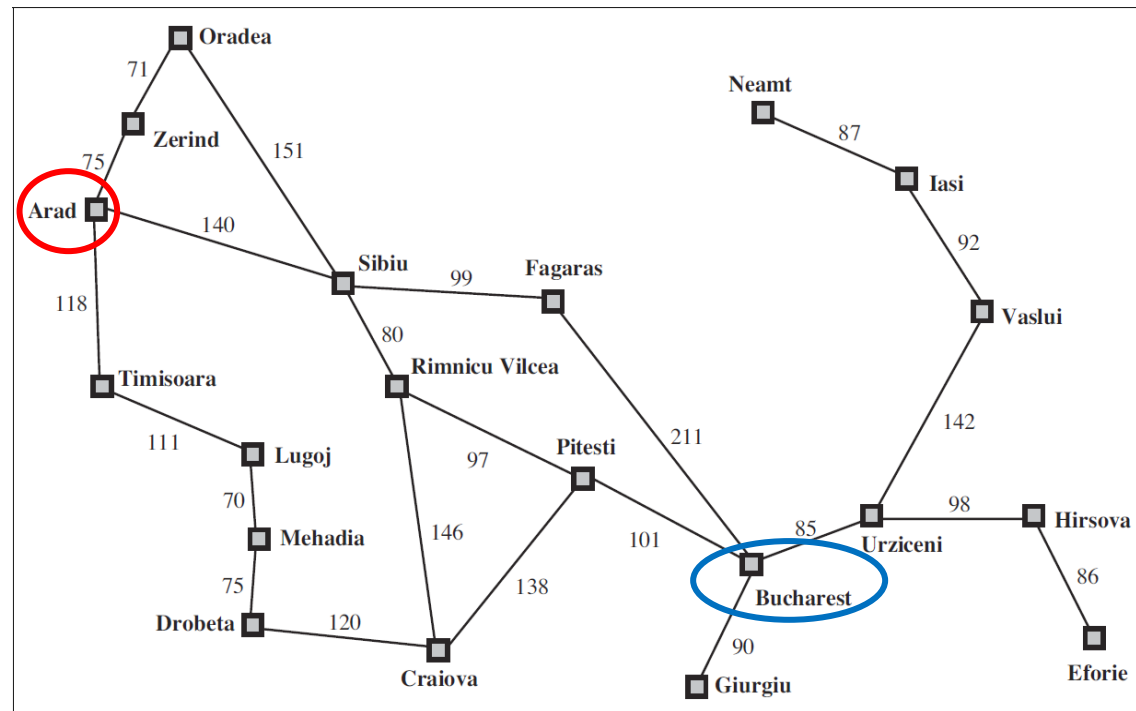
return *action*

Example: Romania Touring

- On holiday in Romania; currently in Arad with a car
- Non-refundable ticket to fly out of Bucharest tomorrow
- **Formulate goal (perf. evaluation):**
 - be in Bucharest before the flight

Road Map of Romania

- **Formulate problem:**
 - states: various cities
 - actions: drive between cities
- **Search:**
 - a sequence of cities



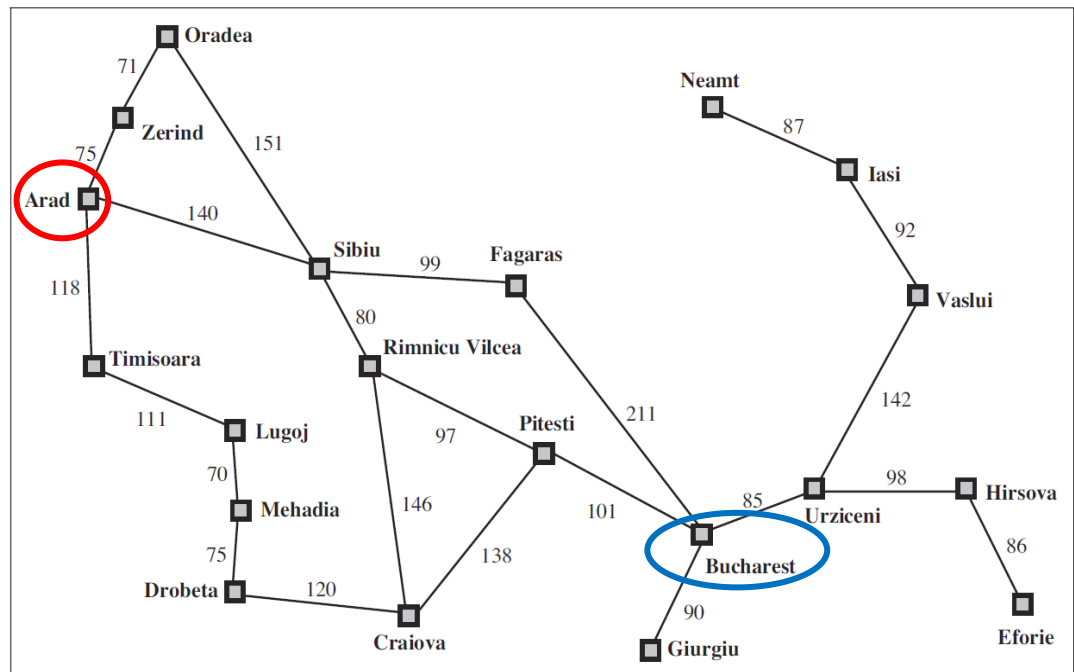
A Well-defined Problem

- A problem can be defined formally by five components
 - Initial state
 - The agent starts in
 - Actions
 - A description of the possible actions available to the agent
 - E.g., Given a state s , $\text{ACTION}(s)$ returns a set of actions that can be executed in s
 - Transition model (successor function)
 - A description of what each action does (returns successor)
 - E.g., $\text{RESULT}(s, a)$ returns the state that results from doing a in s
 - Goal test
 - Determines whether a given state is a goal state
 - Can be **explicit** (goal state) or **implicit** (described by some properties)
 - Path cost function
 - Assigns a numeric cost to each path (a sequence of actions)
 - Step cost (cost of each action)
- Initial state, actions and transition model define the state space
 - the set of all states reachable from the initial state by any sequence of actions
 - forms a graph where nodes are states and the links are actions

Problem Formulation – Example: Romania Touring

- **Initial state:** In(Arad)
- **Actions:**
 - If current state is In(Arad), actions = {Go(Sibiu), Go(Timisoara), Go(Zerind)}
- **Transition model:**
 - e.g., Results(In(Arad), Go(Sibiu)) = In(Sibiu)
- **Goal test:**
 - explicit, e.g. In(Bucharest)
- **Path cost function**
 - e.g., distance traveled

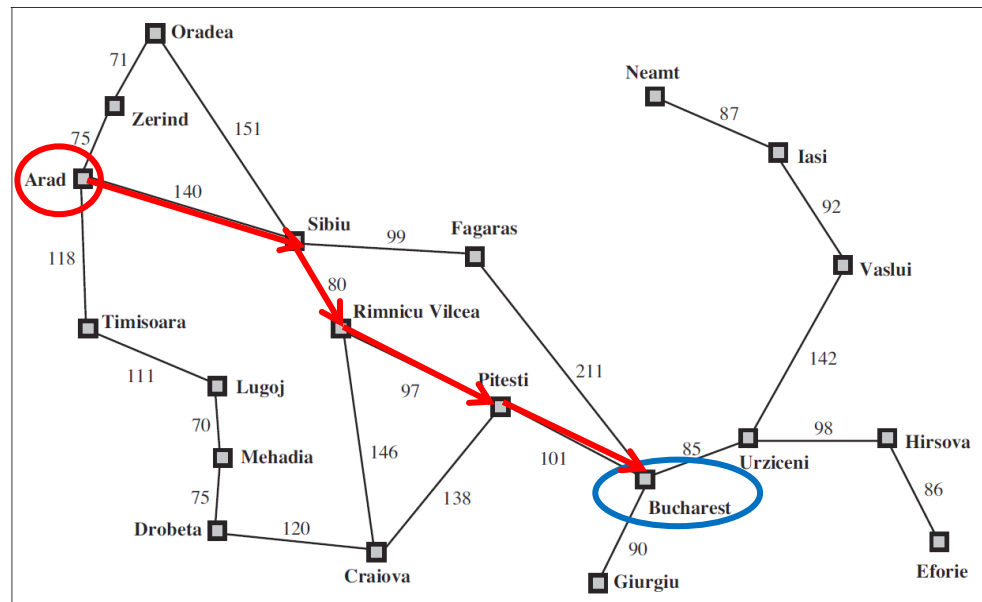
Road Map of Romania



Solution

- A sequence of actions forms a path
- Solution
 - A path from the initial state to a goal state
 - A problem may have more than one solutions
- Optimal Solution
 - Has the lowest path cost among all solutions
 - E.g., Arad, Sibiu, Rimnicu Vilcea, Bucharest

Road Map of Romania



Representation of States and Actions

- **States:**
 - The state of the world includes so many things, e.g., weather, radio, scenery
 - What information is necessary to encode to solving the goal?
 - The size of a problem is usually described in terms of the number of states that are possible.
- **Actions:**
 - The number of actions/operators depends on the representation used in describing a state.
 - E.g., turn left VS turn left 30 degree
- **Good representation for both states and actions can greatly simplify a problem!**

Problem Abstraction

- Real world is complex and has more details
- Abstraction
 - the process to take out the irrelevant information
 - leave the most essential parts to the description of the states
- What's the appropriate level of abstraction?
 - the abstraction is valid, if we can expand it into a solution in a more detailed world
 - the abstraction is useful, if carrying out the actions in the solution is easier than the original problem
 - remove as much detail as possible while retaining validity and usefulness
- Note that every problem formulation involves abstractions

Problem-Solving Agents: Problem Formulation


- **Decide**
 - States?
 - Initial State?
 - Actions?
 - Transition model?
 - Goal test?
 - Path cost?

Review: initial state, actions and transition model define the state space

- the set of all states reachable from the initial state by any sequence of actions
- forms a graph where nodes are states and the links are actions


Example: Vacuum-Cleaner World

- **States**
 - agent location and dirt location
 - number of states: $n * 2^n$ (n squares)
- **Initial state**
 - any state
- **Actions**
 - Left, Right, and Suck [simplified: left out No-op]
- **Transition model**
 - the actions have their **expected** effects
 - complete state space, see next page
- **Goal test**
 - whether both squares are clean
- **Path cost**
 - each step costs 1
 - # of steps in the path



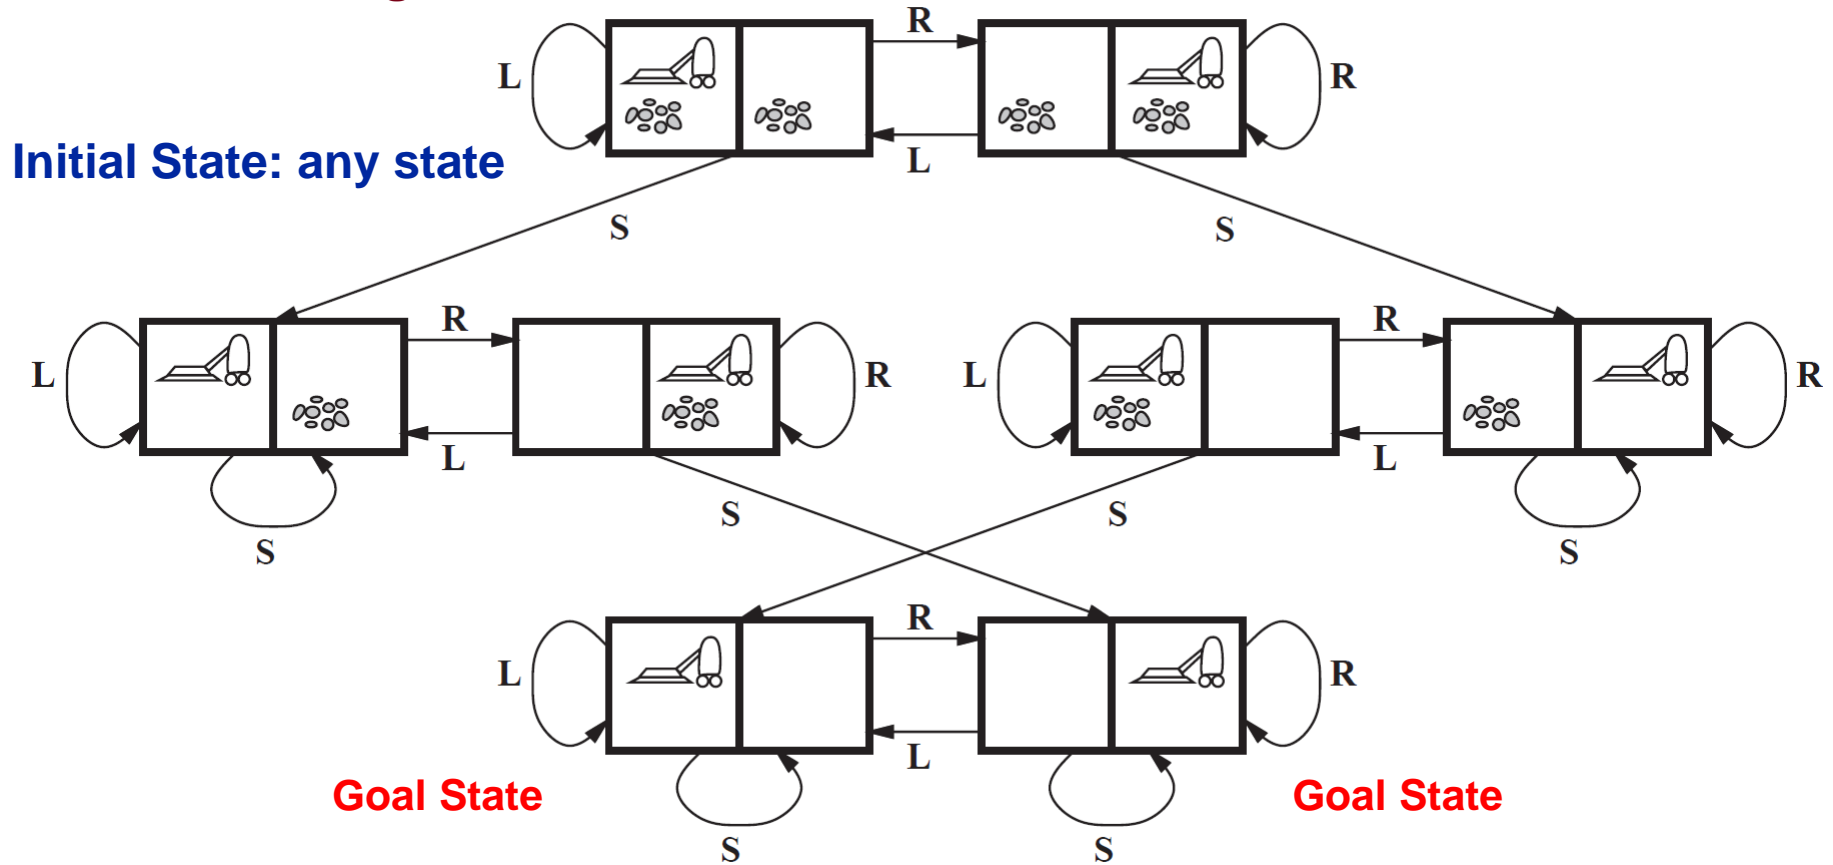
Example: Vacuum-Cleaner World

- The state space graph for the vacuum world.
 - Each node represents a state and each edge denotes an action
- L = Left, R = Right, S = Suck.



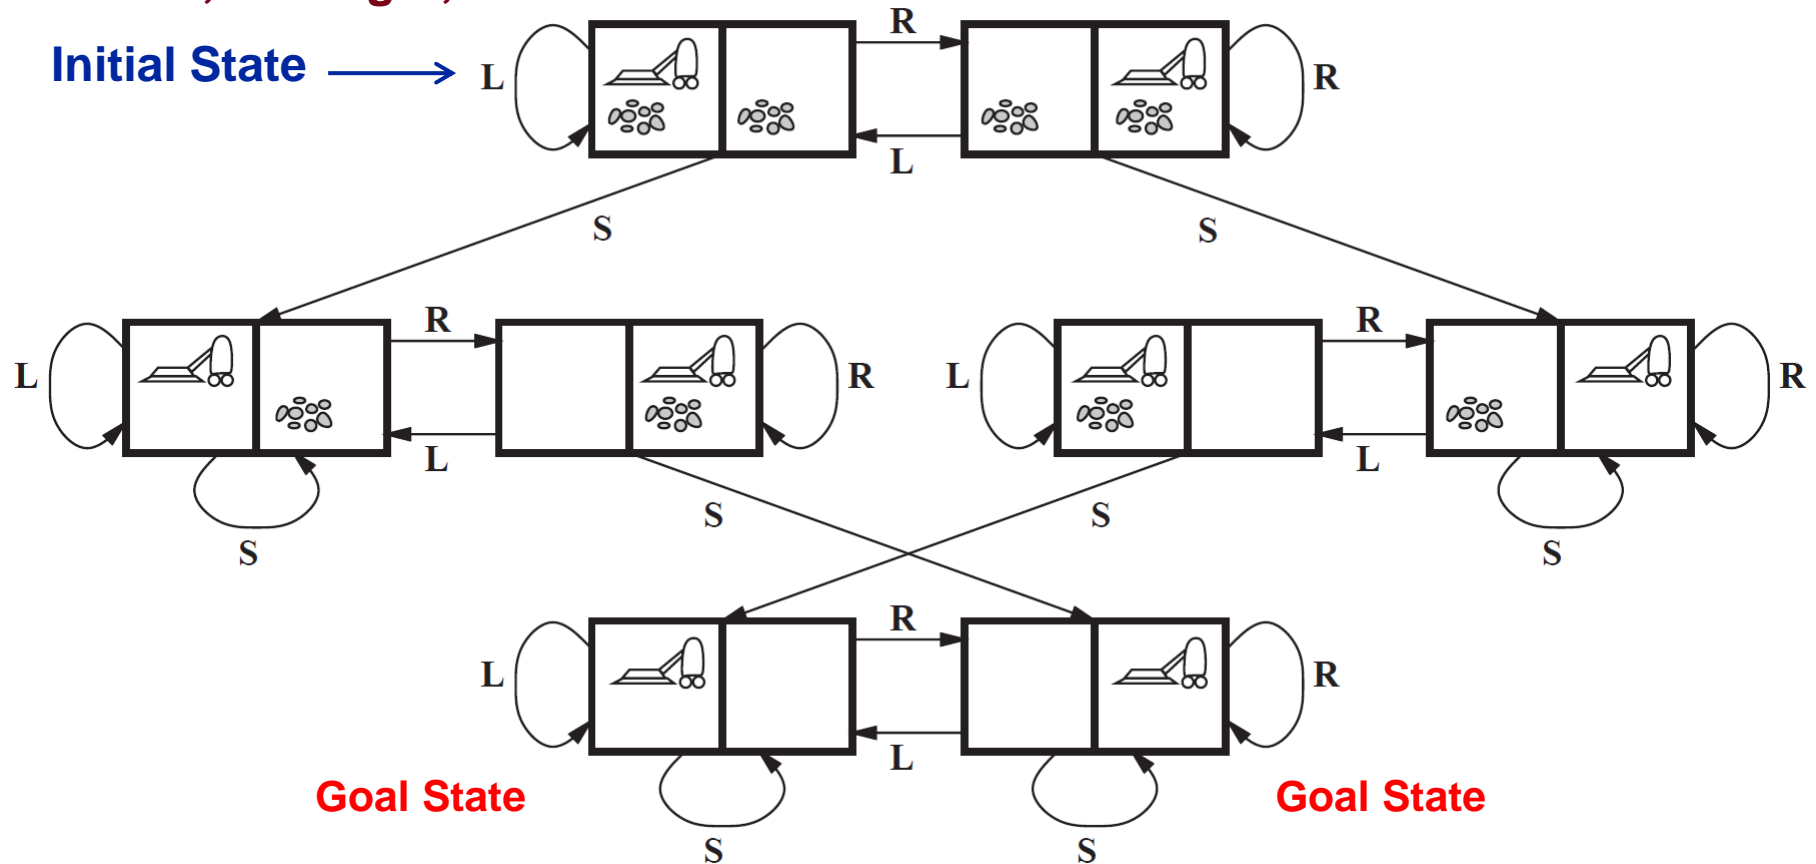
Example: Vacuum-Cleaner World

- The state space graph for the vacuum world.
 - Each node represents a state and each edge denotes an action
- L = Left, R = Right, S = Suck.




Example: Vacuum-Cleaner World

- The state space graph for the vacuum world.
 - Each node represents a state and each edge denotes an action
- L = Left, R = Right, S = Suck.



Example: Vacuum-Cleaner World

- The state space graph for the vacuum world.
 - Each node represents a state and each edge denotes an action
- L = Left, R = Right, S = Suck.



Solution: S \rightarrow R \rightarrow S

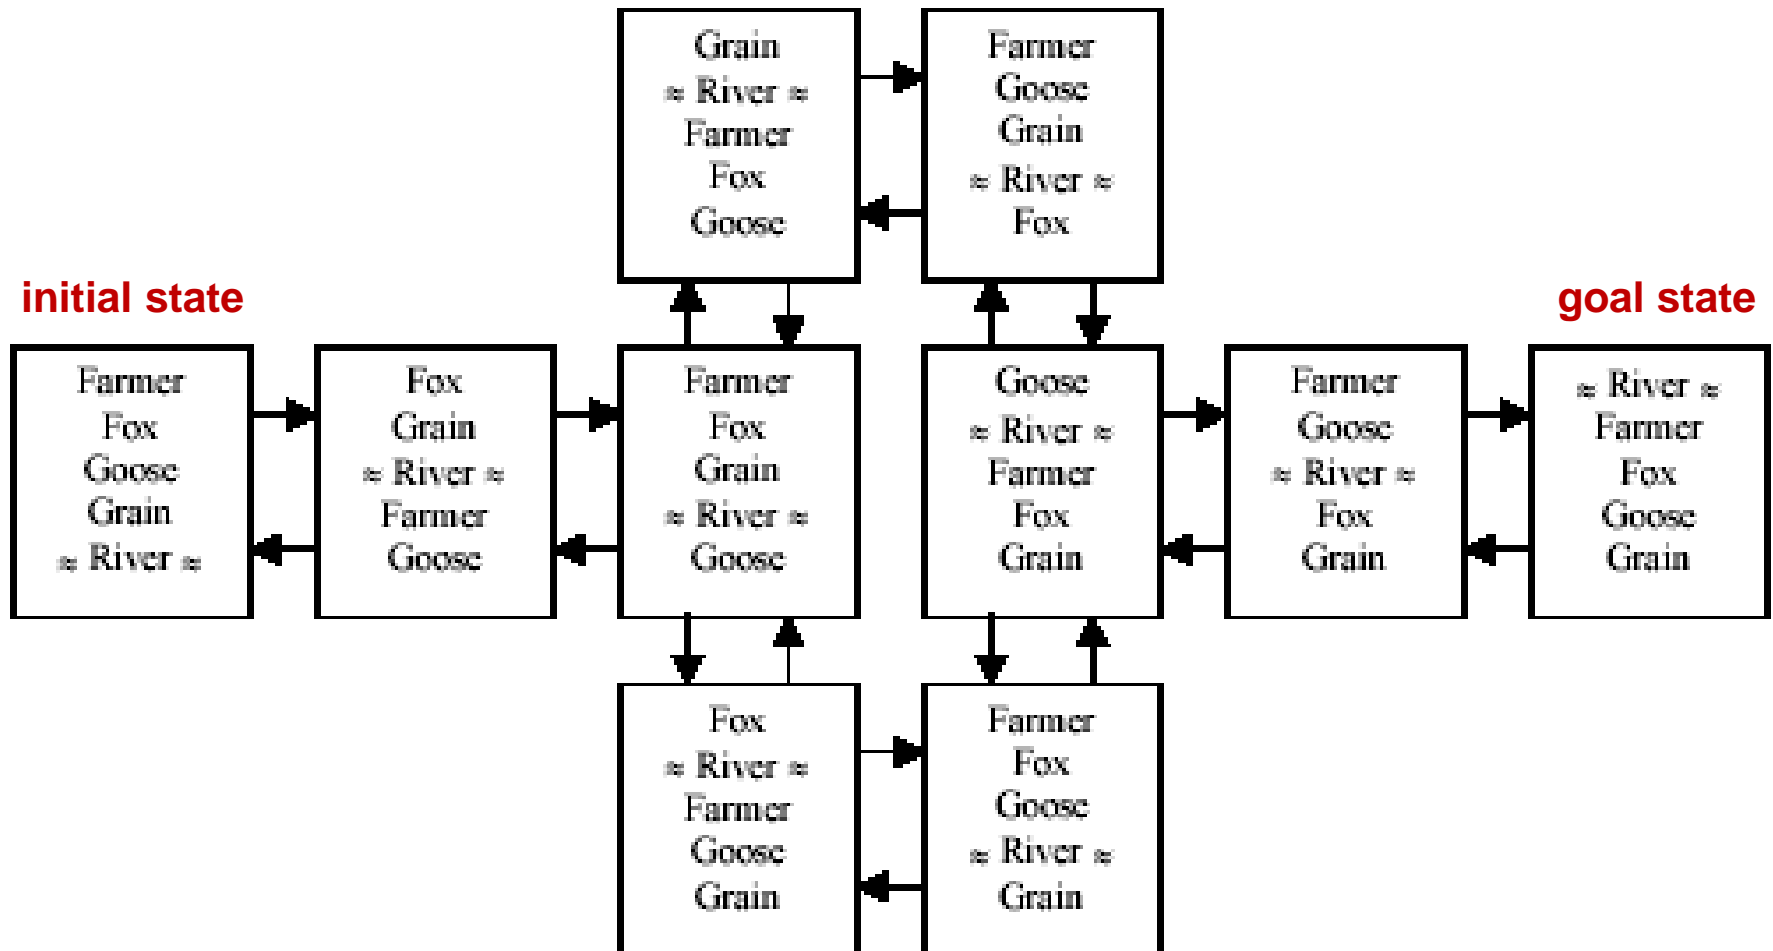
Example: Farmer Problem

- **Good Representations to a Problem are the Key to Good Problem Solving**
- **Example**
 - **Items: Farmer, Fox, Goose, and Grain**
 - **Farmer wants to cross the river with all items by driving a boat**
 - **Fox will eat goose**
 - **Goose will eat grain**
 - **Boat will take max of two and one is the farmer**
- **States?**
- **Initial State?**
- **Actions?**
- **Transition model?**
- **Goal test?**
- **Path cost?**

Problem Formulation for the Farmer Problem

- **States**
 - The locations of the farmer, fox, goose, and grain
- **Initial State**
 - Farmer, fox, goose, and grain are at one side of the river
- **Actions**
 - Farmer moves along with one of the three (fox, goose, grain) to the other side of the river
- **Transition model**
 - Returns a resulting state after an action
- **Goal test**
 - Farmer, fox, goose, and grain are at the other side of the river
- **Path cost**
 - Number of crossing

State Space



Exercise: Missionaries and Cannibals Problem

- **Example**
 - **People: 3 missionaries (black triangles) and 3 cannibals (red dots)**
 - **All 6 people want to cross the river by driving a boat**
 - **The missionaries in one place cannot be outnumbered by cannibals (otherwise the cannibals would eat the missionaries)**
 - **Boat will take one or two people**

- **States?**
- **Initial State?**
- **Actions?**
- **Transition model?**
- **Goal test?**
- **Path cost?**




Problem Formulation

- **States**
 - $\langle m, c, b \rangle$ representing the # of missionaries and the # of cannibals, and the # of the boat at the left side of the river
- **Initial State**
 - $\langle 3, 3, 1 \rangle$
- **Actions**
 - Take 1 missionary, 1 cannibal, 2 missionaries, 2 cannibals, or 1 missionary and 1 cannibal across the river
- **Transition model**
 - Returns a resulting state after an action
- **Goal test**
 - $\langle 0, 0, 0 \rangle$
- **Path cost**
 - Number of crossing



State Space



Representation Principle

Once a problem is formulated using an appropriate representation, the problem is almost solved.

Summary

- **Problem-solving Agent**
- **Goal Formulation**
- **Problem Formulation**

What I want you to do

- Review Chapter 3