

Lecture 6

Local Search

Lusi Li

**Department of Computer Science
ODU**

Reading for This Class:
Chapter 4, Russell and Norvig

Review

- **Last Class**
 - Informed Search
 - Greedy Best-first Search
 - A* Search
 - Heuristic Functions
- **This Class**
 - Local Search
 - Hill-Climbing
 - Simulated Annealing
- **Next Class**
 - Global Search
 - Genetic Algorithms

Path Search VS. Local Search

- The search algorithms we discussed before are designed to find a goal state from a start state s
 - the **path** to the goal that constitutes a **solution** to the problem
 - Uninformed search: $g(s)$
 - Informed search: $h(s)$, $g(s)+h(s)$
- In many optimization problems
 - the path to the goal is irrelevant
 - **the goal state itself is the solution**
- Another category of search problem
 - Local Search Problem
 - Never worry about the path
 - **Just want the goal**
 - Examples
 - Integrated-circuit Design
 - Factory-floor layout
 - Automatic programming
 - Telecommunications network optimization

Path Search VS. Local Search

- **Path Search** maintains a search tree to find the path
 - keep paths in memory and remember alternatives
 - can backtrack
- **Local search** uses a **single search path** of solutions, not a search tree
 - start from an initial state
 - at each step consider **the current state**, and try to improve it by moving only to one of its neighbors (not the one in frontier set)
 - Iterative improvement algorithms
 - No frontier set and no backtracking
 - Each state is a solution

Local Search Algorithms


- **Objective (Fitness) Function $f(s)$**
 - All local search problems have an objective function to specify how “good” a state is
 - Each state s has a score $f(s)$
 - The goal is to find the state with the highest (or the lowest) score, or a reasonably high (or low) score
- **General Procedure**
 - Keep only a single (complete) state in memory
 - Generate only the neighbours of that state
 - Keep one of the neighbors and discard others
- **Two strategies for choosing the state to visit next**
 - Hill climbing
 - Simulated annealing
- **Then, an extension to multiple current states**
 - Genetic algorithms

Local Search Algorithms

- **Two key advantages**
 - Very little memory required
 - Often find reasonable solutions in large or infinite state spaces
- **Usage**
 - Pure optimization problem
 - Find or approximate the best state according to some objective function
 - Optimal if the space to be searched is convex

1-D State Space Landscape

- **Global maximum**
 - Find the highest peak
- **Global minimum**
 - Find the lowest valley



A one-dimensional state-space landscape in which elevation corresponds to the objective function.

Hill Climbing Algorithm

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*

if neighbor.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

- **Idea: start from some state *s***
 - move to a neighbor *t* with a better score *f(t)*. Repeat.
- **Properties:**
 - Terminate when no neighbor has better value
 - Does not look ahead beyond the immediate neighbors of the current state
 - Choose randomly among the set of best successors, if there is more than one
 - Do not backtrack, since it doesn't remember where it's been
 - Required data structure: **the current state and the *f(s)***
 - a.k.a. greedy local search

Hill Climbing Algorithm

- **Q1: What's a neighbor?**
- **Q2: Pick which neighbor?**
- **Q3: What if no neighbor is better than the current state?**

Hill Climbing Algorithm


- **Q1: What's a neighbor?**
 - You have to define that!
 - The neighborhood of a state is the set of neighbors.
 - Also called 'move set'
 - Similar to successor function
- **Q2: Pick which neighbor?**
 - The best one (greedy) based on objective function values
- **Q3: What if no neighbor is better than the current state?**
 - Stop

Hill-climbing: 8-Queens problem

- Put all 8 queens on the 8 x 8 board with no two queens attacking each other, i.e, no two queens can share the same row, column, or diagonal.
- Complete state formulation:
 - **State:**
 - **Neighbor states:**
 - **Fitness function f :**

Constraints:

1. Each row must contain exactly one queen.
2. Each column must contain exactly one queen.
3. No two queens should be in the same row, column, or diagonal.




Hill-climbing: 8-Queens problem

- Put all 8 queens on the 8 x 8 board with no two queens on the same row, column, or diagonal
- Complete state formulation:
 - **State**: positions of the 8 Queens one per column
 - **Neighbor states**: generated by moving one queen to a different square in the same column
 - **Fitness function f**: number of pairs of queens that are attacking each other
 - Note that we want a state s with the lowest score $f(s) = 0$
 - Low or high should be obvious from context.









Constraints:

1. Each row must contain exactly one queen.
2. Each column must contain exactly one queen.
3. No two queens should be in the same row, column, or diagonal.



8-Queens problem: fitness values of neighborhood

$$f(s) = 3 + 4 + 2 + 3 + 2 + 2 + 1 + 0 = 17$$

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14		13	16	13	16
	14	17	15		14	16	16
17		16	18	15		15	
18	14		15	15	14		16
14	14	13	17	12	14	12	18

An 8-queens state with $f(s) = 17$.


It also shows the value of f for each possible successor obtained by moving a queen within its column, with the best one having $f = 12$.

The best moves are marked.

Hill-climbing algorithms typically choose **randomly** among the set of best successors if there is more than one.

Fig. 1 An 8-queens state s

8-Queens problem: Local minimum



An 8-queens state with $f(t) = 1$.

It is a local optimum because every move leads to a larger f .

Fig. 2 An 8-queens state t

Performance of 8-Queens Problem

- **8-queens statistics:**
 - State space of size ≈ 17 million
 - Starting from random state, steepest-ascent hill climbing solves 14% of problem instances and gets stuck for 86% of problem instances
 - It takes 4 steps on average when it succeeds, 3 when it gets stuck
 - When **sideways** moves are allowed, performance improve
 - Sideways moves: if no uphill moves, allow moving to a state with the same value as the current one.
 - E.g., 100 consecutive sideways moves, 14% \rightarrow 94%

Analysis of Hill-Climbing

- **Continually moves uphill**
 - increasing value of the evaluation function
 - (or “downhill” decreasing value of the cost function)
 - gradient descent search is a variation that moves downhill
- **Very simple strategy with low space requirements**
 - stores only the state and its evaluation, no search tree
- **Problems**
 - local maxima
 - plateau
 - ridges

Analysis of Hill-Climbing

- **Problems**

- local maxima



- the peak is higher than all its neighbors, but not the global maximum
 - algorithm can't go higher, but is not at a satisfactory solution

- plateau

- area where the evaluation function is flat
 - the best neighborhood has the same value as the current state

- ridges

- sequence of local maxima difficult for greedy algorithms to navigate
 - search may oscillate slowly



Further Variants of Hill Climbing

- **Sideways moves:**
 - if no uphill moves, allow moving to a state with the same value as the current one (escape shoulders)
- **Stochastic hill-climbing:**
 - **selection** among the available uphill moves is done to be “less” greedy
 - the better, the more likely
- **First-choice hill-climbing:**
 - successors are generated **randomly**, one at a time, until one that is better than the current state is found
 - if better, take the move
 - deal with large neighborhoods
- **Random-restart hill climbing:**
 - conducts a series of hill-climbing searches from randomly generated initial states, until a goal is found.
 - “If at first you don’t succeed, try, try again.”


Simulated Annealing

- **Escape from local optima**
 - by accepting, with a **probability** that decreases during the search, also moves that are worse than the current solution (going “downhill”)
- **Inspired by the process of annealing of solids in metallurgy:**
 - annealing: harden metals and glass by heating them to a high temperature and then gradually cooling them, thus allowing the material to reach a low-energy crystalline state
 - at the start, make lots of moves and then gradually slow down

Simulated Annealing: Intuition

- Minimization problem
- Imagine a state space landscape on table
- Let ping-pong ball from random point → local minimum
- Shake table → ball tends to find different minimum
- Shake hard at first (high temperature) but gradually reduce intensity (lower temperature)

Ball on terrain example – SA vs Greedy Algorithms



Upon a large no. of iterations, SA converges to the solution.

Simulated Annealing

1. Pick an initial state s
2. Randomly pick t in $\text{neighbors}(s)$
3. IF $f(t)$ better THEN accept $s \leftarrow t$.
4. ELSE /* t is worse than s */
5. accept $s \leftarrow t$ **with a small probability**
6. GOTO 2 until bored.

Q: How to choose the small probability?

Simulated Annealing

1. Pick an initial state s
2. Randomly pick t in $\text{neighbors}(s)$
3. IF $f(t)$ better THEN accept $s \leftarrow t$.
4. ELSE /* t is worse than s */
5. accept $s \leftarrow t$ **with a small probability**
6. GOTO 2 until bored.

Q: How to choose the small probability?

- idea 1: $p = 0.1$
- idea 2: p decreases with time
- idea 3: p decreases with time, also as the difference between $f(t)$ and $f(s)$ increases

Simulated Annealing

1. Pick initial state s
2. Randomly pick t in neighbors(s)
3. IF $f(t)$ better THEN accept $s \leftarrow t$.
4. ELSE /* t is worse than s */
5. accept $s \leftarrow t$ **with a small probability**
6. GOTO 2 until bored.


Q: How to choose the small probability?

- idea 1: $p = 0.1$
- idea 2: p decreases with time
- **idea 3:** p decreases with time, also as the difference between $f(t)$ and $f(s)$ increases

Simulated Annealing

- $\Delta E = f(t) - f(s)$
- If $f(t)$ is better than $f(s)$, always accept t
- Otherwise, accept t with probability

$$p = e^{\frac{\Delta E}{T}}$$




- i.e., if $r < p$ ($r \in [0, 1]$ is a uniform random number), accept t
- where T is a temperature parameter that 'cools' (anneals) over time, e.g. $T \leftarrow T * 0.9$
 - High T allows more worse moves
 - Low T results in few or no bad moves
- If the difference (formally known as energy difference) $|f(t) - f(s)|$ is large, the probability is small.

Simulated Annealing

- Acceptance criterion and cooling schedule

if ($\Delta \geq 0$) accept

else if ($random < e^{\Delta / Temp}$) accept, else reject /* $0 \leq random \leq 1$ */



Initially temperature is very high (most bad moves accepted)

Temp slowly goes to 0, with multiple moves attempted at each temperature

Final runs with $temp=0$ (always reject bad moves) greedily “quench” the system

Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

for $t = 1$ **to** ∞ **do**

$T \leftarrow \text{schedule}(t)$ // **T is the current temperature, which is monotonically decreasing with t**

if $T = 0$ **then return** *current* // **halt when temperature = 0**

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow \text{next.VALUE} - \text{current.VALUE}$ // **If positive, next is better than current.**

if $\Delta E > 0$ **then** *current* \leftarrow *next* // **Otherwise, next is worse than current.**

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

// **as $T \rightarrow 0$, $p \rightarrow 0$; as $\Delta E \rightarrow -\infty$, $p \rightarrow 0$**

Summary

- **Local Search**
- **Hill-climbing**
- **Simulated Annealing**

What I want you to do

- Review Chapter 4