

Lecture 13

Games II

Lusi Li

**Department of Computer Science
ODU**

Reading for This Class:
Chapter 5, Russell and Norvig


Review

- **Last Class**
 - Games
 - Adversarial Search
- **This Class**
 - Pruning a Game Tree
- **Next**
 - Midterm Exam (02/29-03/01)


Review: Optimal Strategy

- MAX moves first and they take turns moving until game is over
- Assumption: both players play optimally.
 - given a choice, MAX prefers to move to a state of maximum value
 - whereas MIN prefers a state of minimum value
- Given a game tree, each node will get a MiniMax value.
- High value is good for MAX but bad for MIN.

MIN




MAX




Review: Optimal Strategy

- MAX moves first and they take turns moving until game is over
- Assumption: both players play optimally.
 - given a choice, MAX prefers to move to a state of maximum value
 - whereas MIN prefers a state of minimum value
- Given a game tree, each node will get a MiniMax value.
- High value is good for MAX but bad for MIN.

MIN



MAX



Review: MiniMax Value

MiniMax(s) =

if **Terminal-Test(s)** then **Utility(s)**

if **Player(s) = MAX** then

 max of **MiniMax(Result(s, a))** for **a** in **Actions(s)**

if **Player(s) = MIN** then

 min of **MiniMax(Result(s, a))** for **a** in **Actions(s)**

Review: Minimax Algorithm

function MINIMAX-DECISION(*state*) **returns** *an action* // choose the best move
return $\arg \min_{a \in \text{ACTIONS}(s)} \text{MAX-VALUE}(\text{RESULT}(s, a))$ // for a MIN node

function MINIMAX-DECISION(*state*) **returns** *an action* // choose the best move
return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(s, a))$ // for a MAX node

function MAX-VALUE(*state*) **returns** *a utility value* // calculate the value
if **TERMINAL-TEST**(*state*) **then return** **UTILITY**(*state*) // for a MAX node
 $v \leftarrow -\infty$ // when finding a higher value, update v
for each *a* **in** **ACTIONS**(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$
return *v*

function MIN-VALUE(*state*) **returns** *a utility value* // calculate the value
if **TERMINAL-TEST**(*state*) **then return** **UTILITY**(*state*) // for a MIN node
 $v \leftarrow \infty$ // when finding a lower value, update v
for each *a* **in** **ACTIONS**(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$
return *v*

MiniMax Properties


- Assume all terminal states are at depth m and there are b possible moves at each step
- Minimax explores tree using DFS.
 - recursive implementation
- Time complexity is $O(b^m)$
 - exponential in the number of moves
- Space complexity is $O(bm)$
 - where b is the branching factor, m the maximum depth of the search tree
- Completeness: Yes, if tree is finite
- Optimality: Yes, if MIN is an optimal opponent

Problem with Minimax Search

Problems:

- **Complete search is impractical for most games**
 - **Number of game states is exponential in the number of moves**
- **Solution: Make the same decision without examining every node by pruning**
- **Alpha-beta pruning**
 - **an extension of the minimax approach**
 - **results in the same move as minimax, but with less overhead**
 - **prunes uninteresting parts of the search tree**

Intuition of Alpha-Beta Pruning



$$\begin{aligned}\text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \\ &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\ &= 3.\end{aligned}$$

Therefore, it is possible to compute the **correct** minimax decision **without looking at every node** in the tree.

Alpha-Beta Pruning

- α :
 - has initial value $-\infty$
 - the value of the best (i.e., highest-value) choice we have found so far at any choice point along the path for MAX
 - $\alpha \leftarrow \text{MAX}(\alpha, v)$ updated only when it's MAX's turn
- β
 - has initial value $+\infty$
 - the value of the best (i.e., lowest-value) choice we have found so far at any choice point along the path for MIN.
 - $\beta \leftarrow \text{MIN}(\beta, v)$ updated only when it's MIN's turn

Alpha-Beta Pruning

Key points:

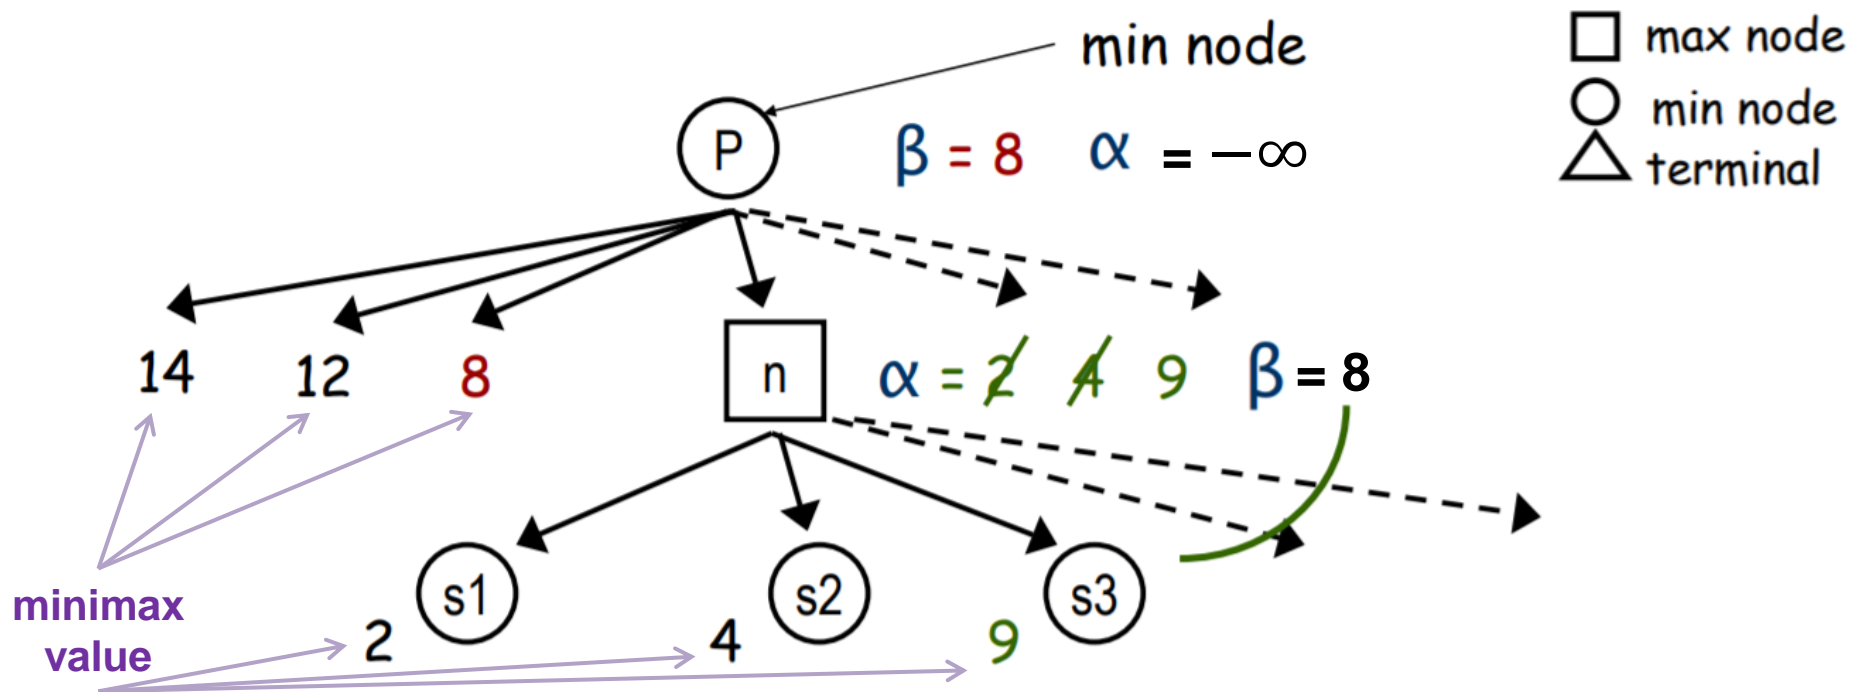
1. Each node has to keep track of 3 values: α , β , v (minimax value).
2. Pruning condition: if $\alpha \geq \beta$ for node n , stop expanding the children of node n and return its current v
3. MAX will update only α values and MIN player will update only β values. Both of them will update v values.
4. Return v values to **parent** nodes of the game tree
5. Pass α and β values to **child** nodes.

Two types of pruning (cuts):

- pruning of max nodes (α -cuts)
- pruning of min nodes (β -cuts)

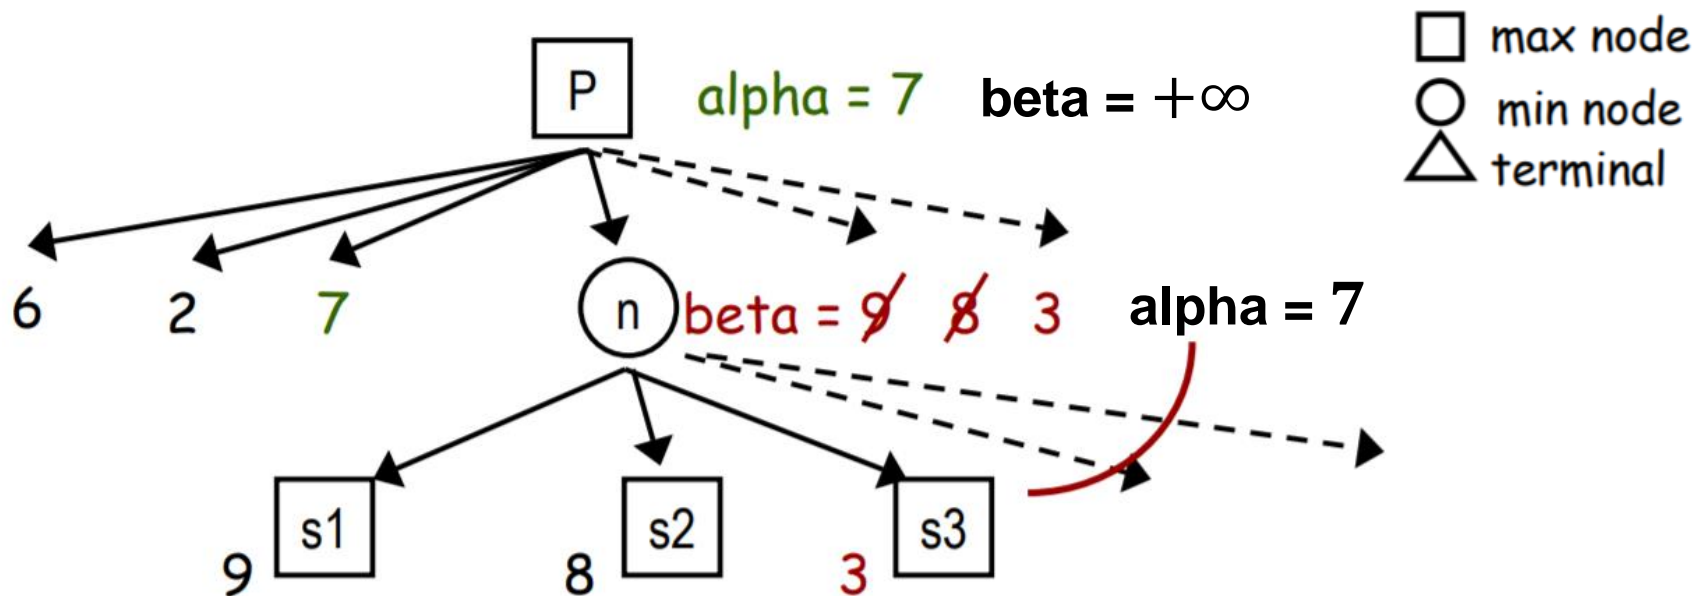
Alpha-Beta Pruning: α -cut

- If $\alpha \geq \beta$ for node n , stop expanding the children of node n
- α -cut: cutting some of the child nodes of a max node
 - MIN will never choose to move from n 's parent (P) to n since it would choose one of n 's lower valued siblings first.



Alpha-Beta Pruning: β -cut

- If $\alpha \geq \beta$ for node n , stop expanding the children of node n
- β -cut: cutting some of the child nodes of a min node
 - Max will never choose to move from n 's parent (P) to n since it would choose one of n 's higher value siblings first.



Alpha-Beta Algorithm

function ALPHA-BETA-SEARCH($state$) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(state, -\infty, +\infty)$
 return the *action* in $\text{ACTIONS}(state)$ with value v

// choose the best move
// for a MAX node

function MAX-VALUE($state, \alpha, \beta$) **returns** a utility value
 if $\text{TERMINAL-TEST}(state)$ **then return** $\text{UTILITY}(state)$
 $v \leftarrow -\infty$
 for each a **in** $\text{ACTIONS}(state)$ **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \geq \beta$ **then return** v // prune remaining children of this max node
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 return v // return value of best child

// calculate v and α
// for a MAX node

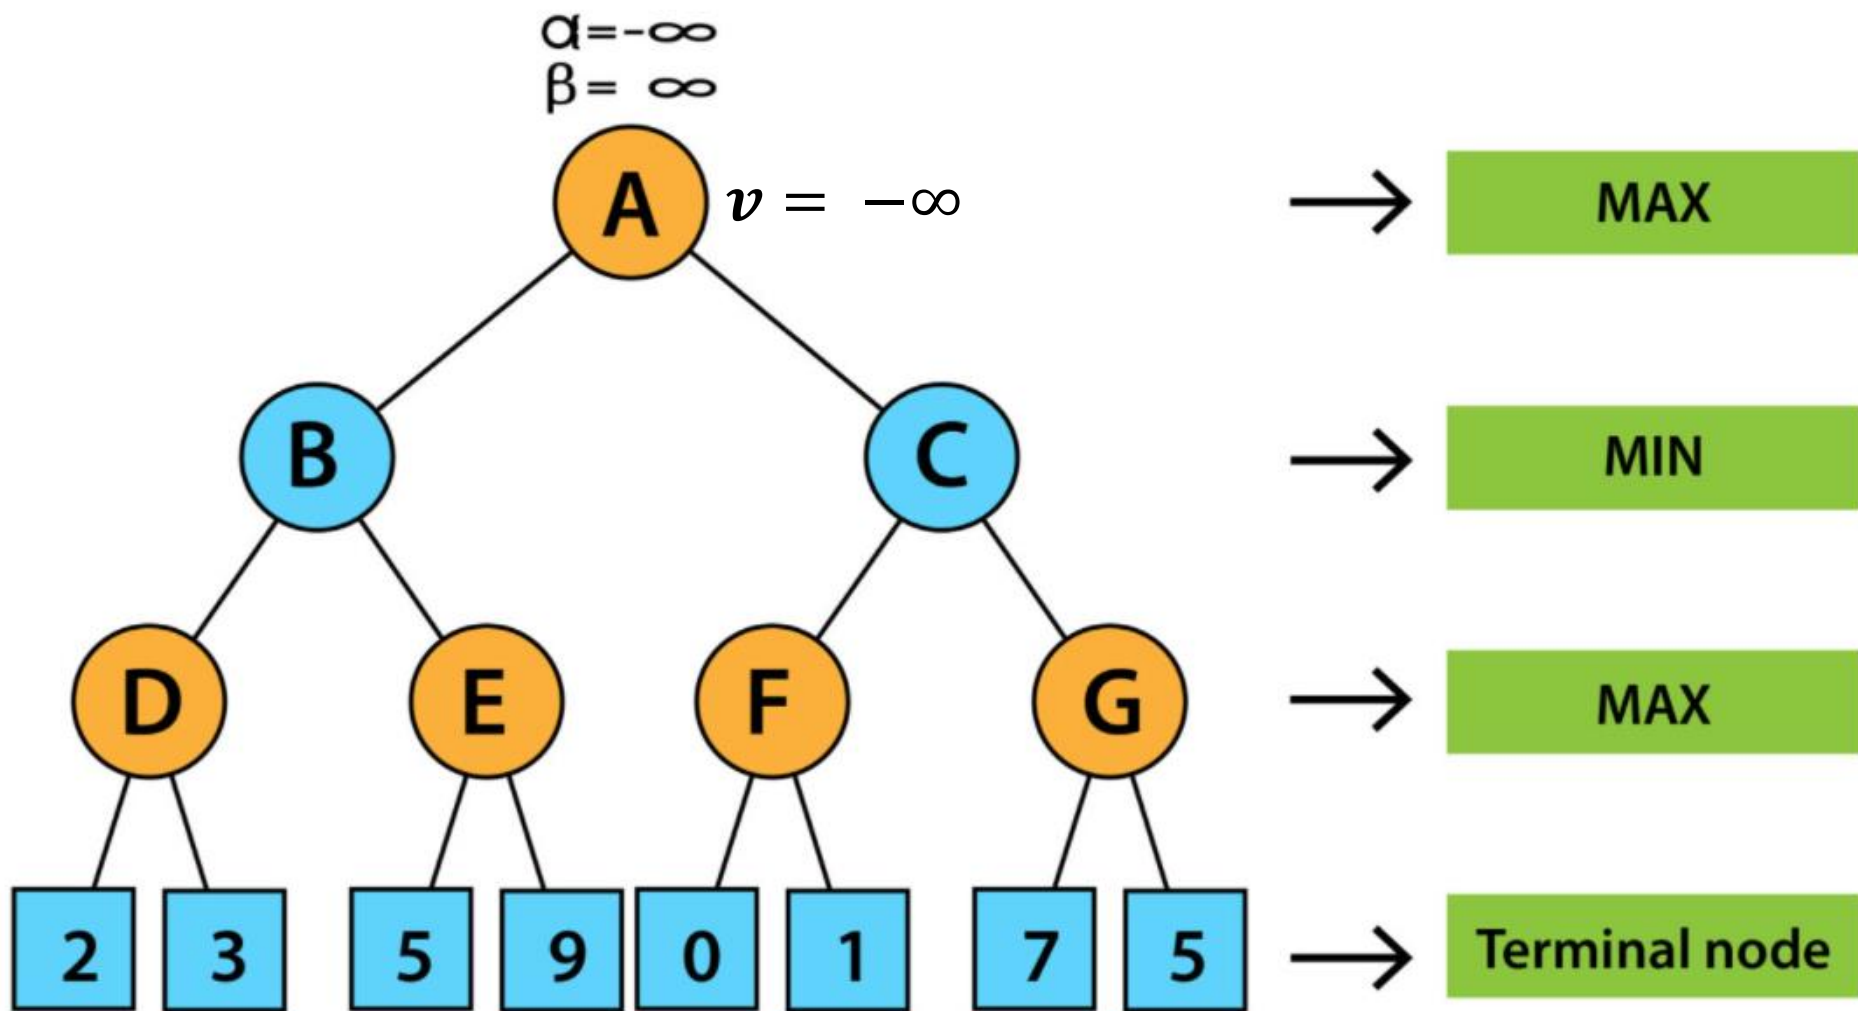
function MIN-VALUE($state, \alpha, \beta$) **returns** a utility value
 if $\text{TERMINAL-TEST}(state)$ **then return** $\text{UTILITY}(state)$
 $v \leftarrow +\infty$
 for each a **in** $\text{ACTIONS}(state)$ **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \leq \alpha$ **then return** v
 $\beta \leftarrow \text{MIN}(\beta, v)$
 return v

// calculate v and β
// for a MIN node

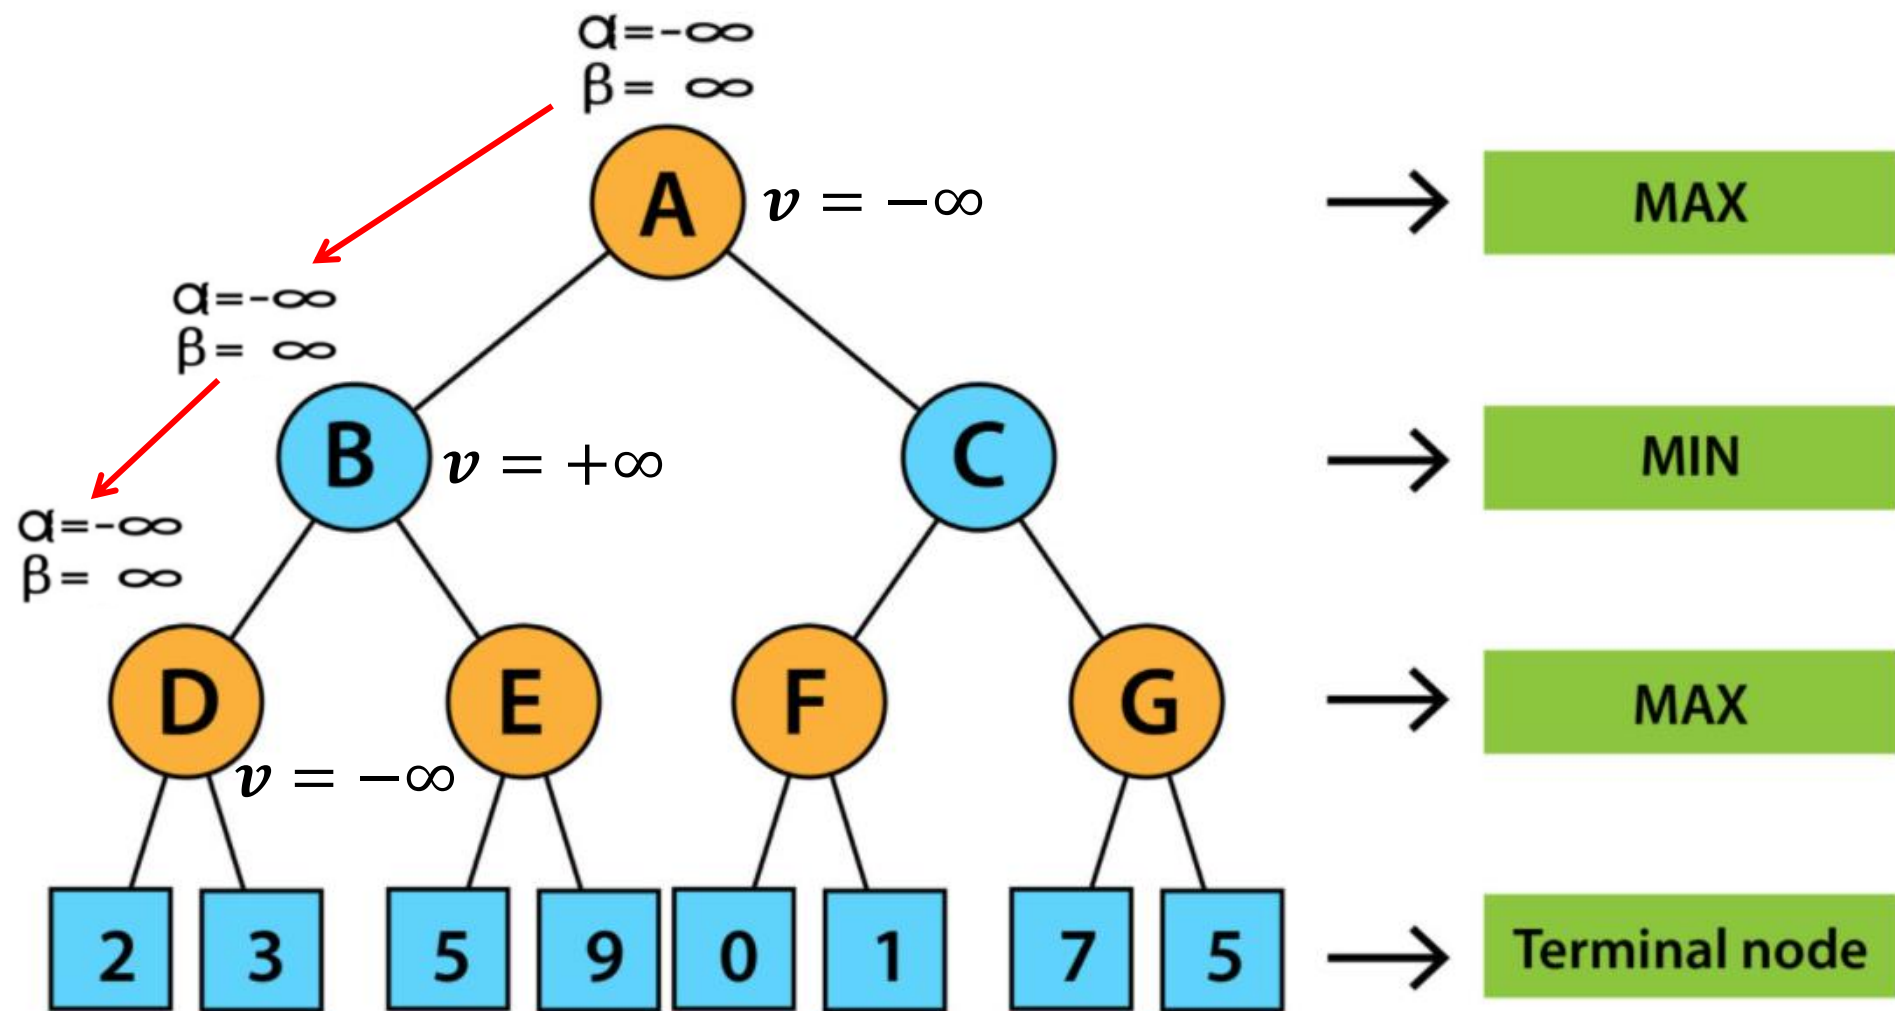
return v

Alpha-Beta Pruning Example

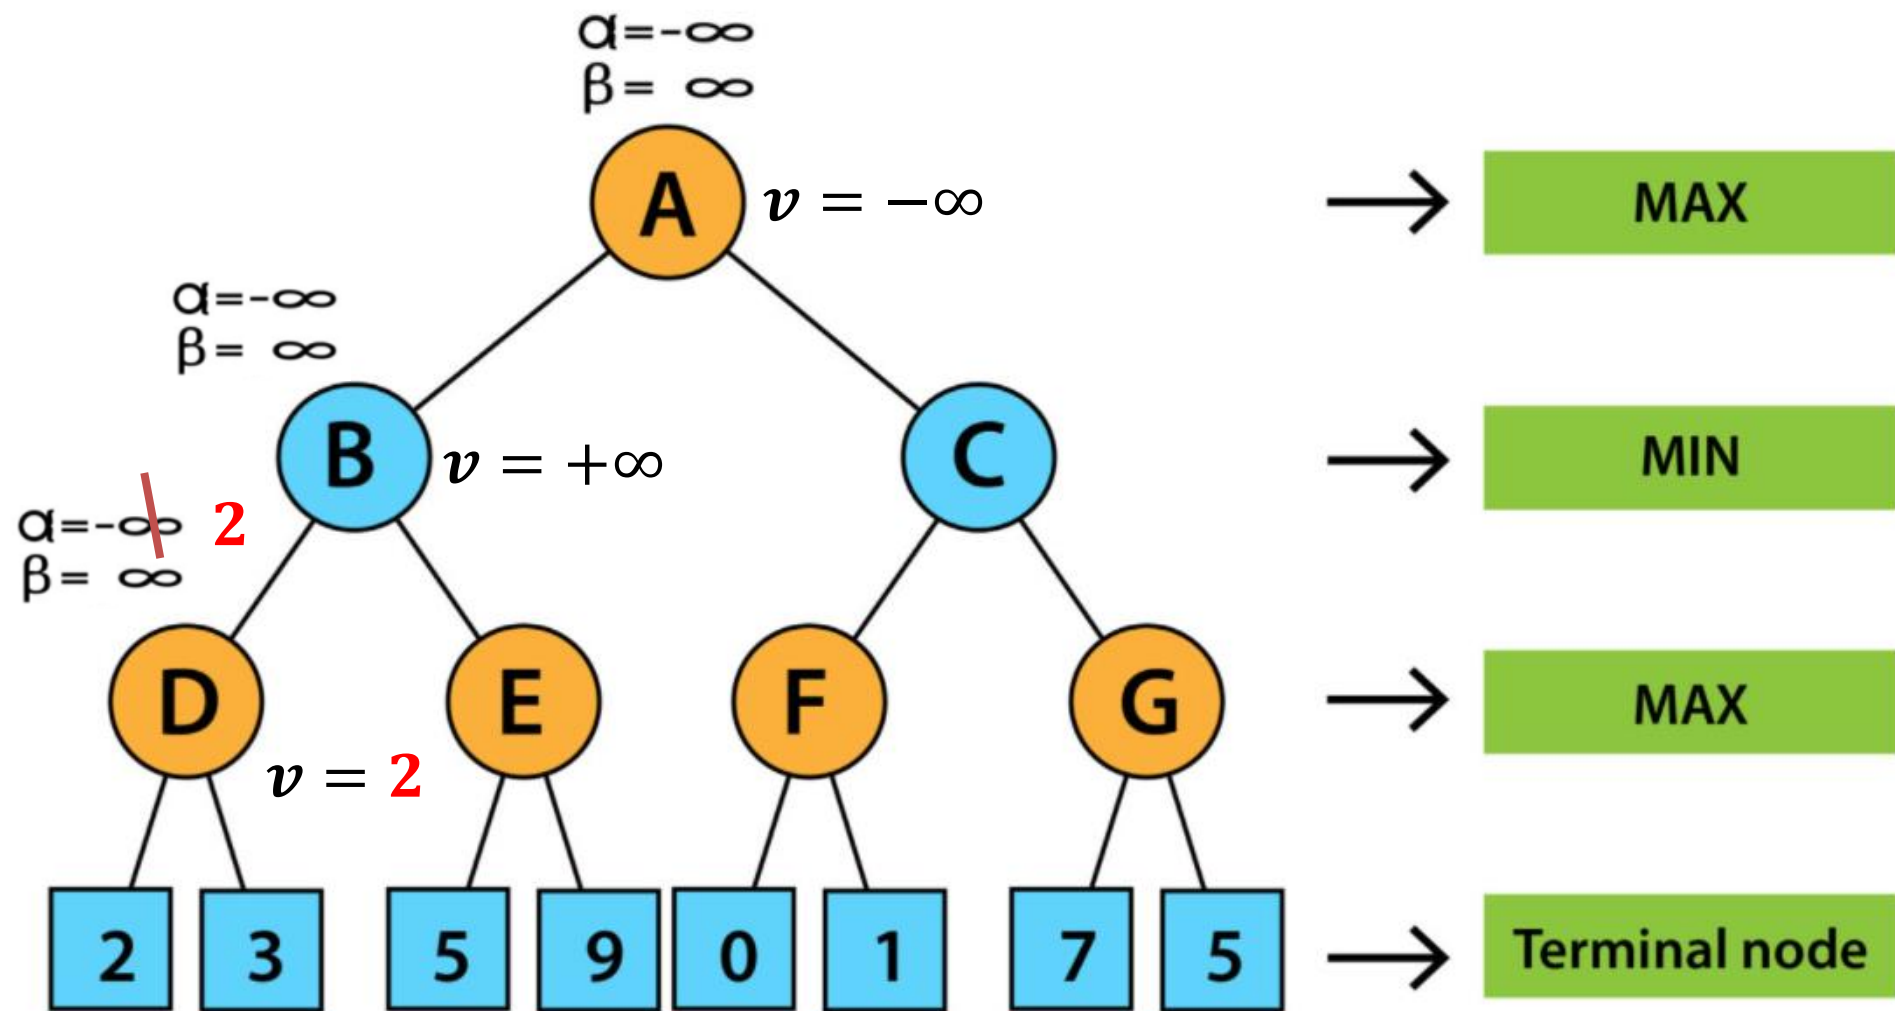
ALPHA-BETA-SEARCH (A) **returns** *an action*



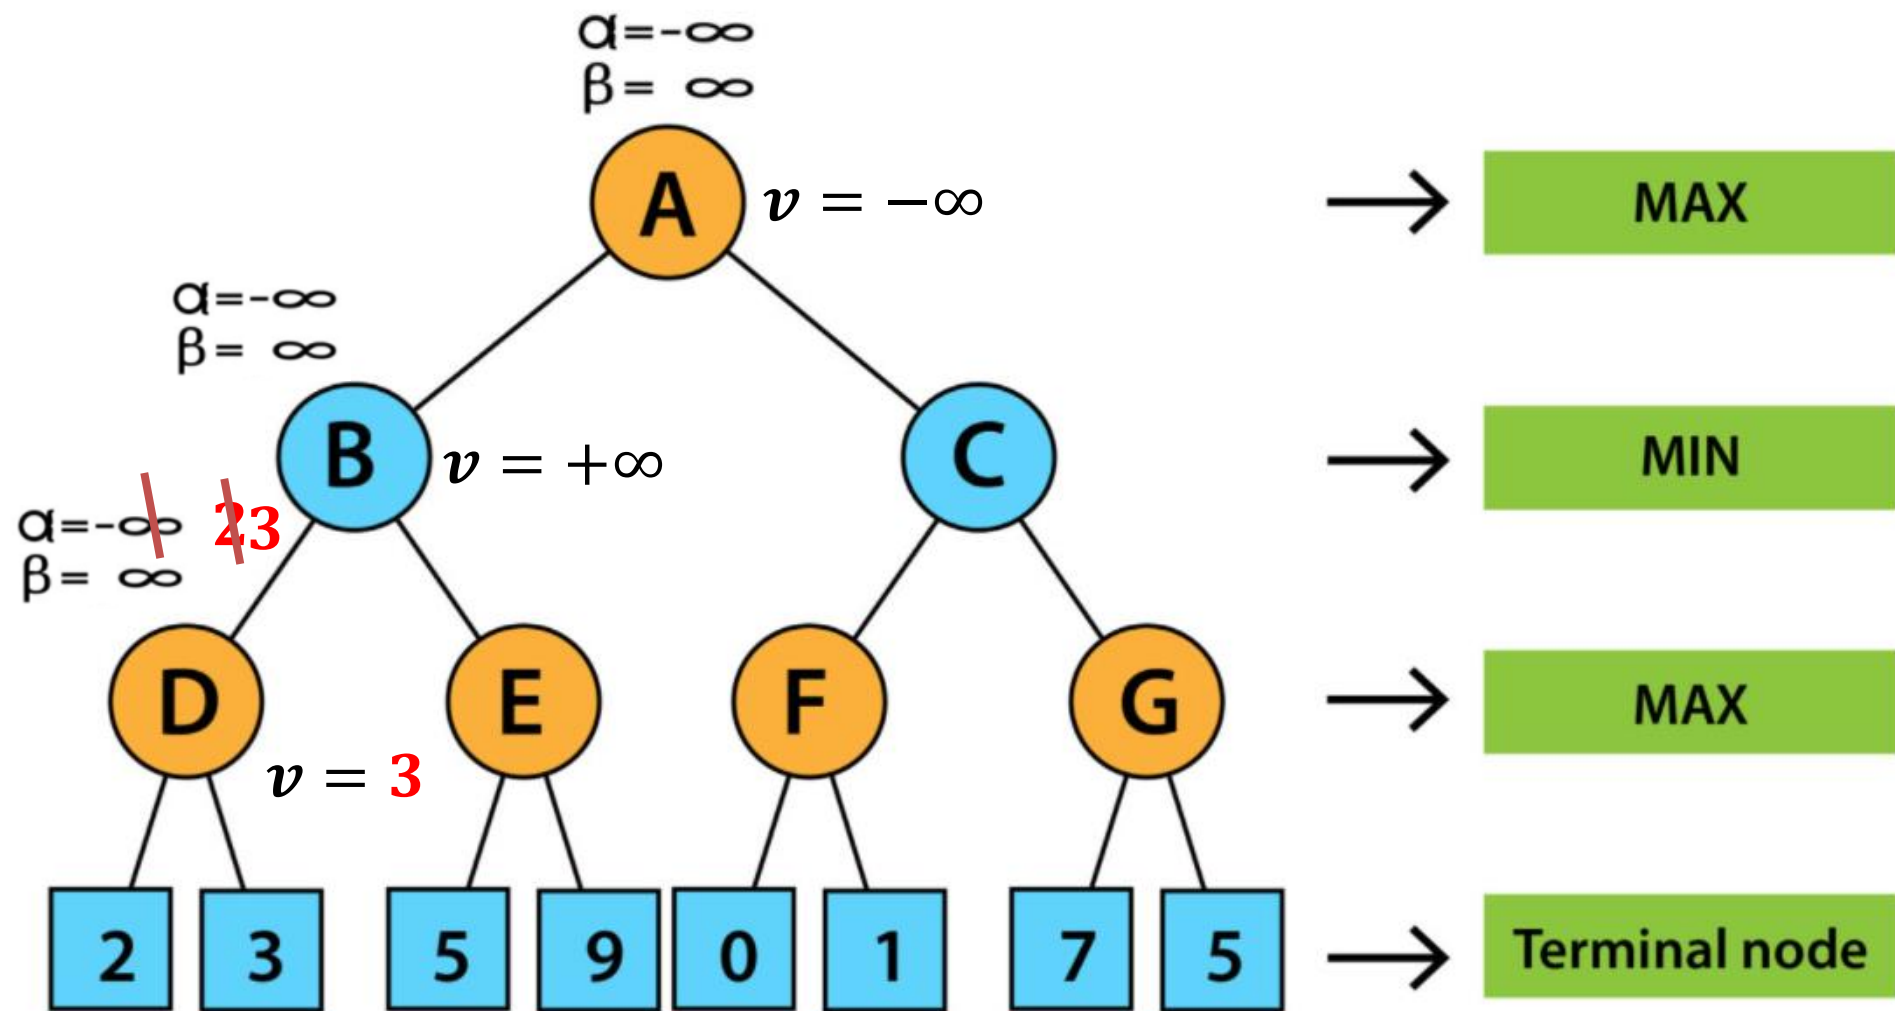
Alpha-Beta Pruning Example



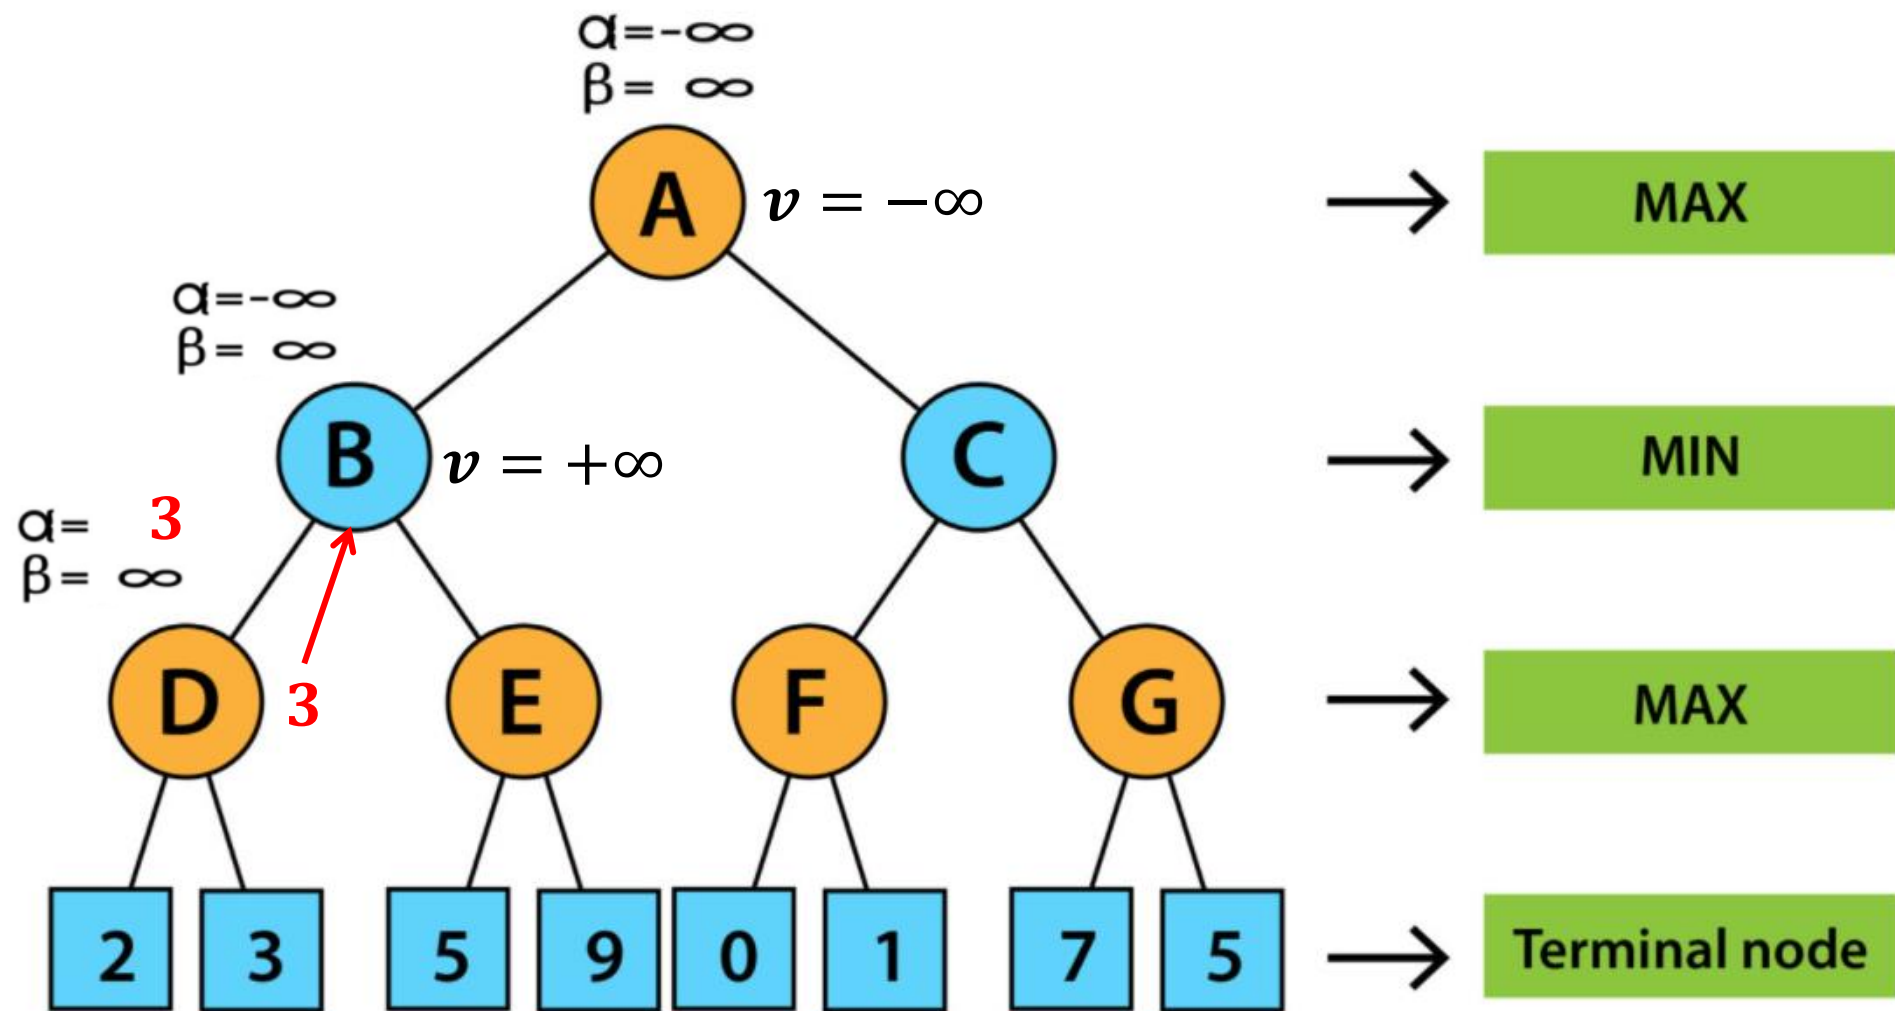
Alpha-Beta Pruning Example



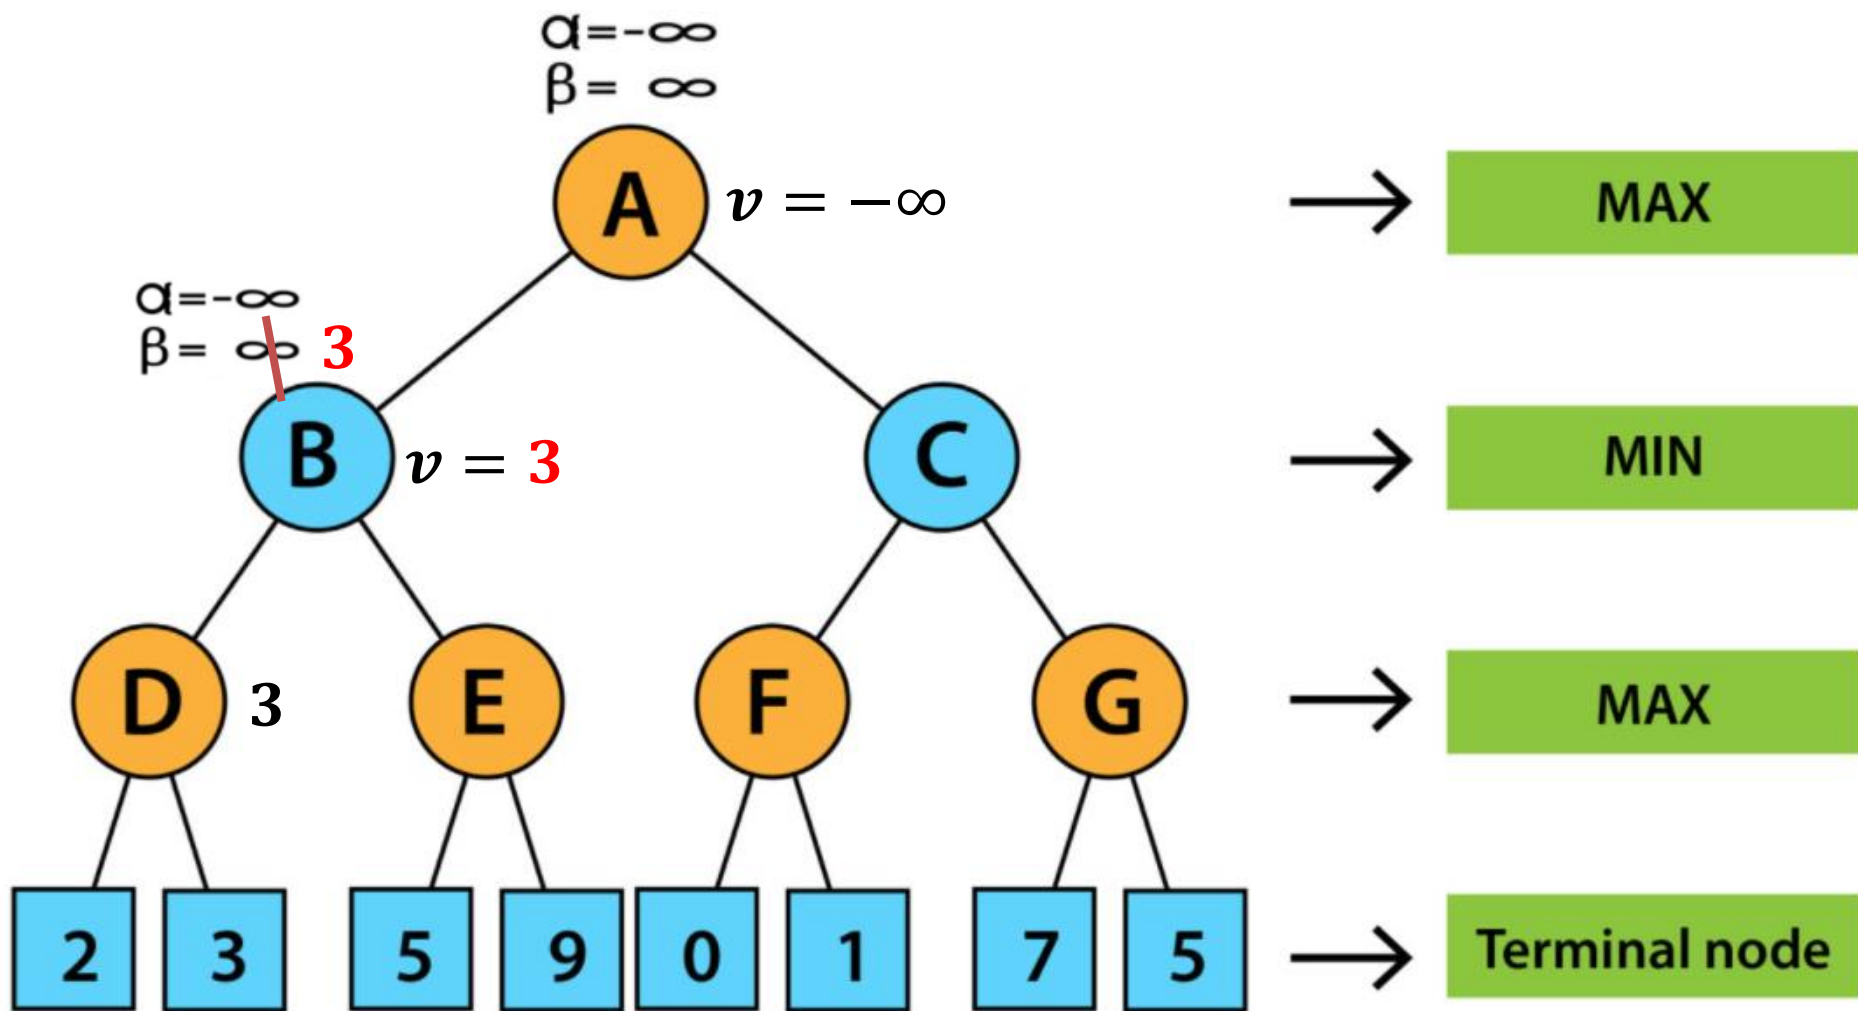
Alpha-Beta Pruning Example



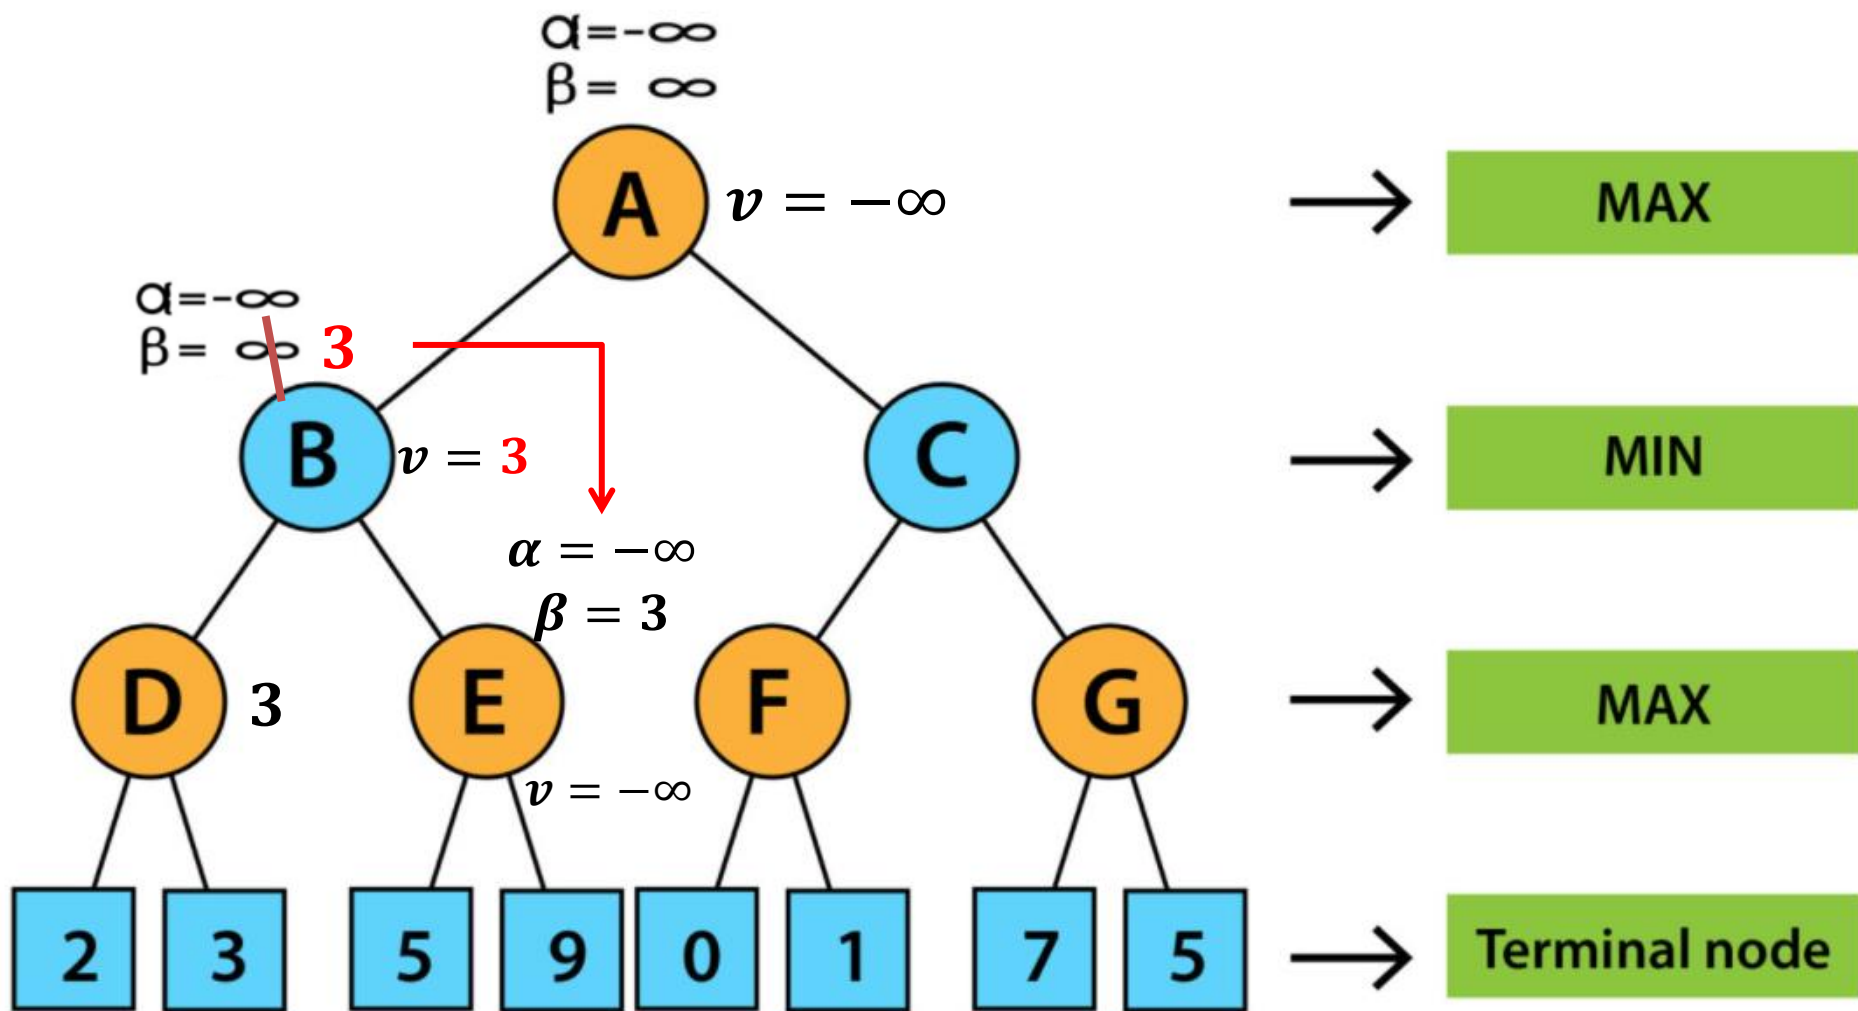
Alpha-Beta Pruning Example



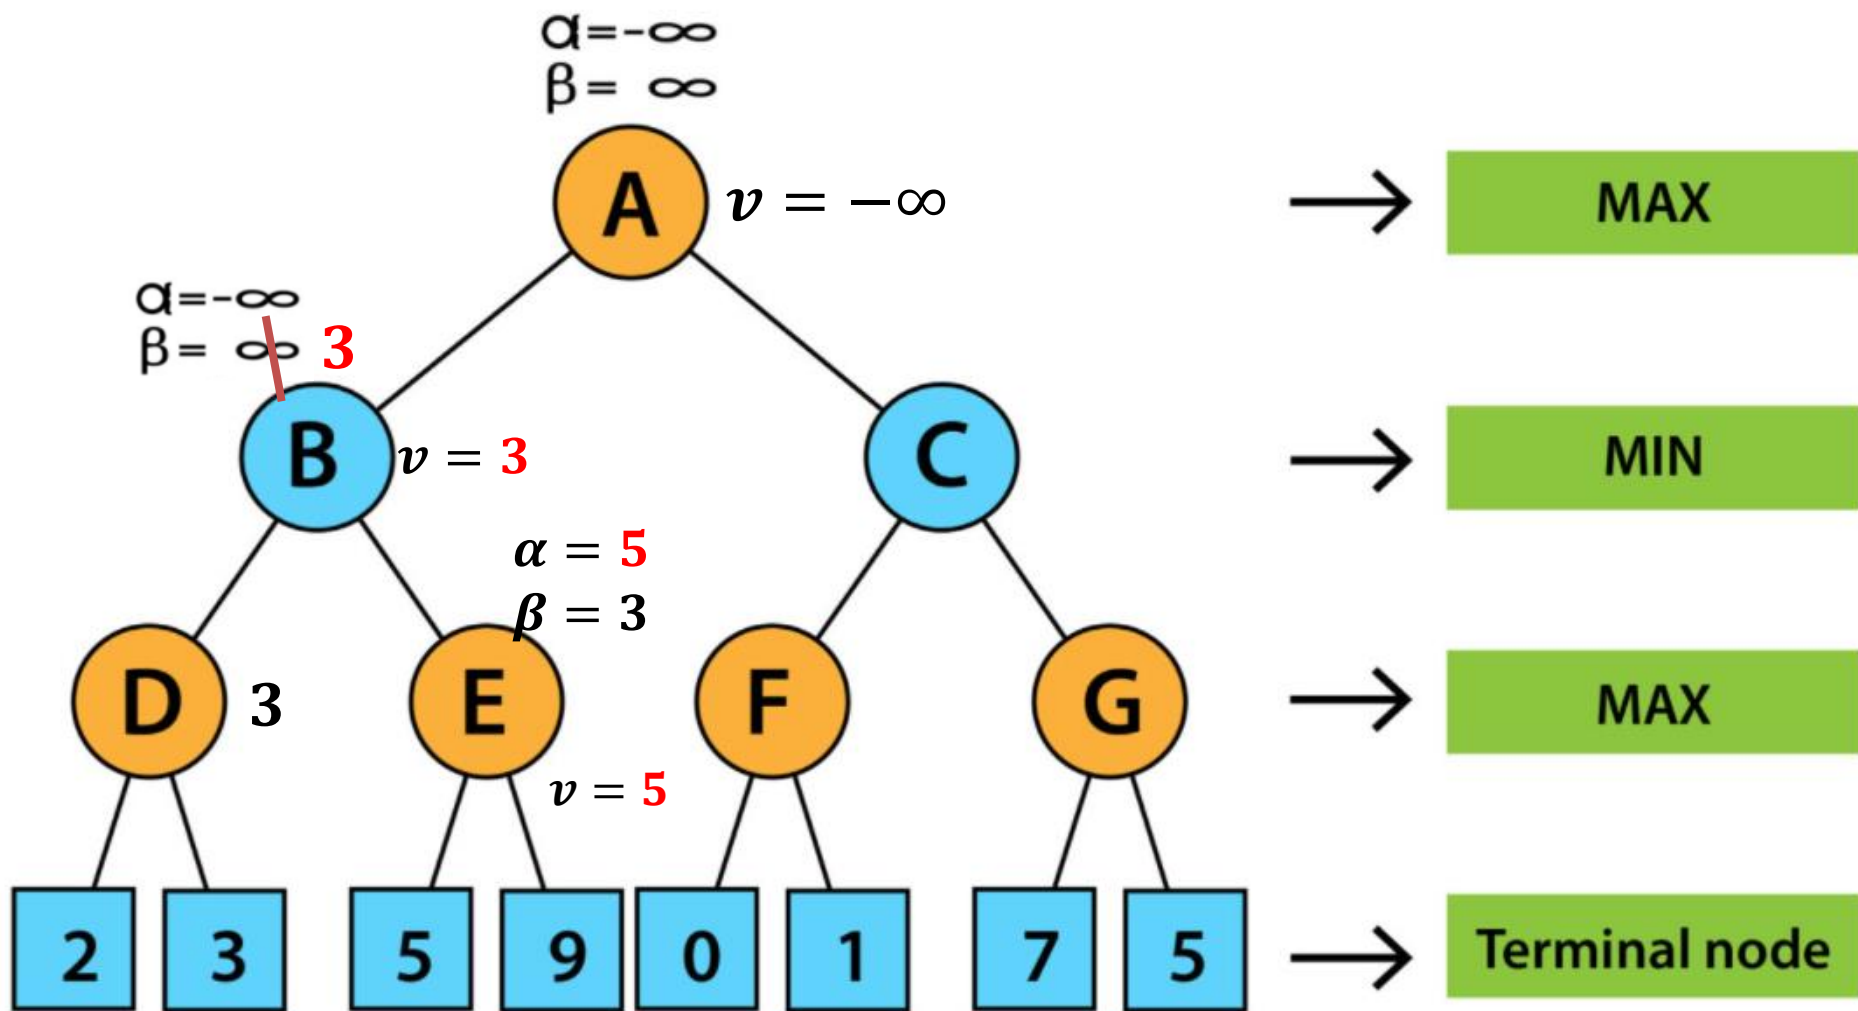
Alpha-Beta Pruning Example



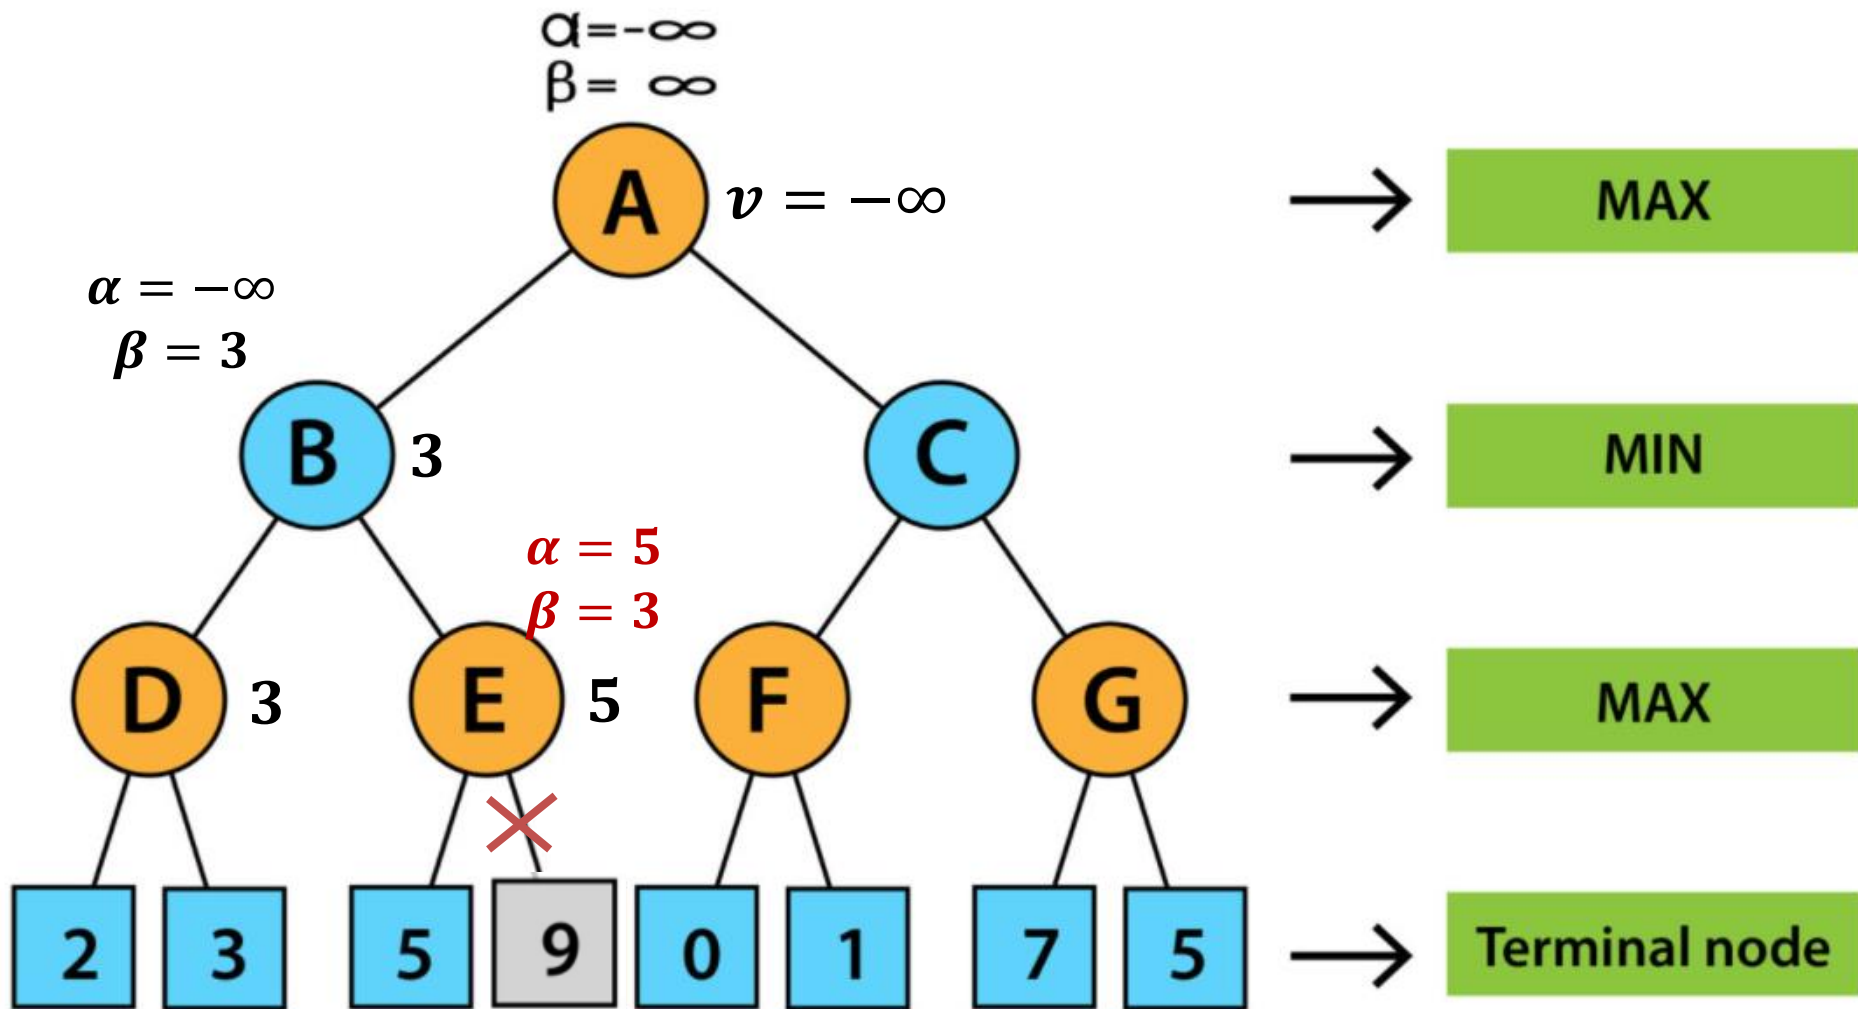
Alpha-Beta Pruning Example



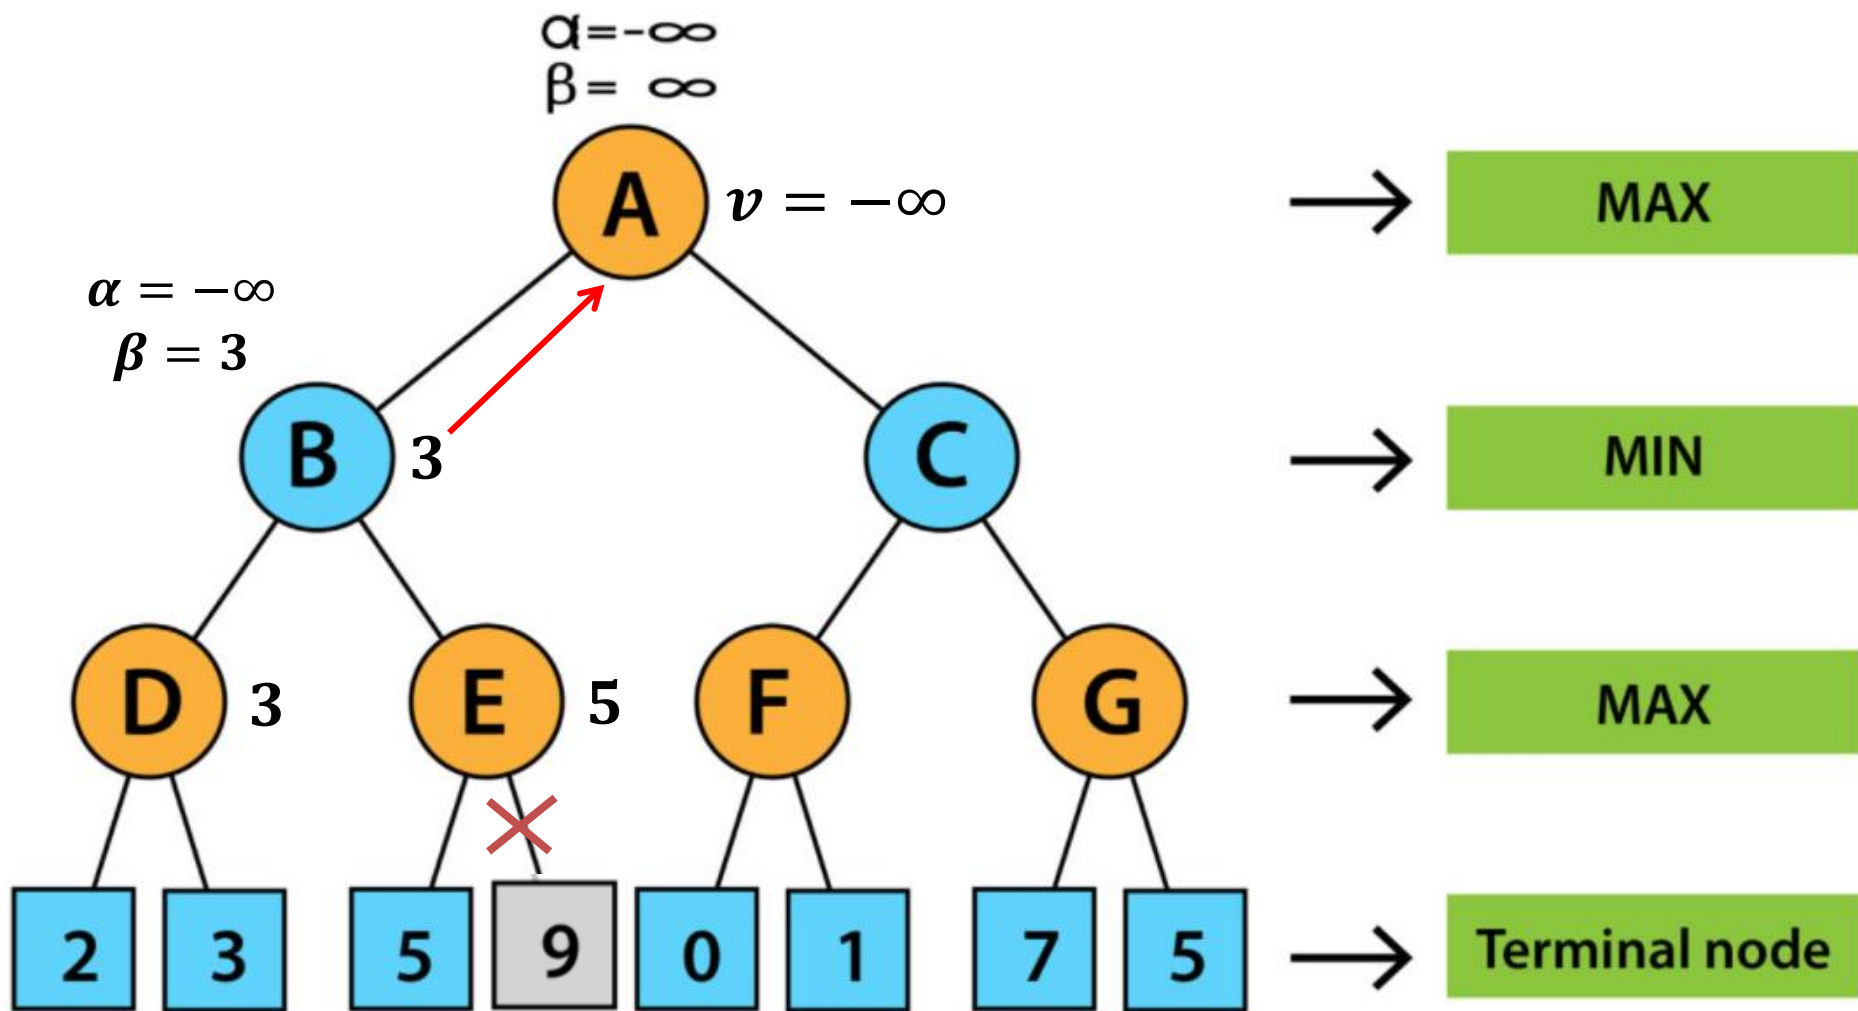
Alpha-Beta Pruning Example



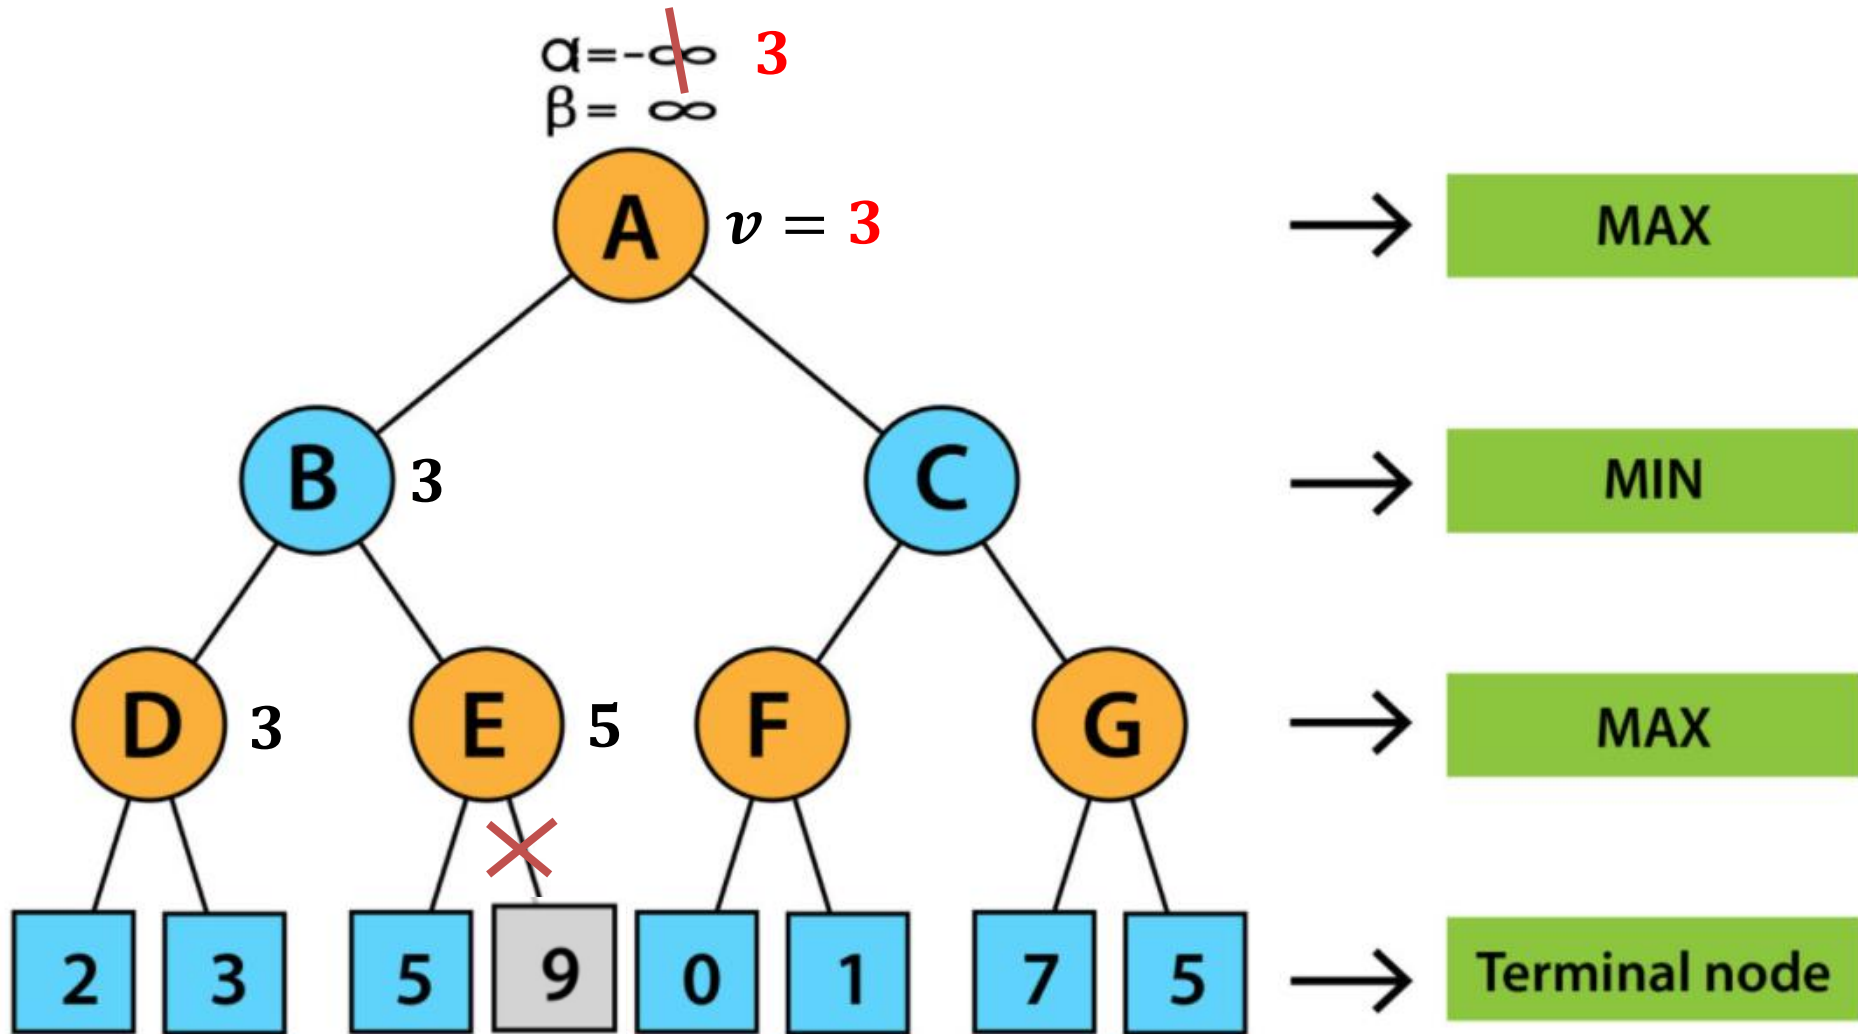
Alpha-Beta Pruning Example



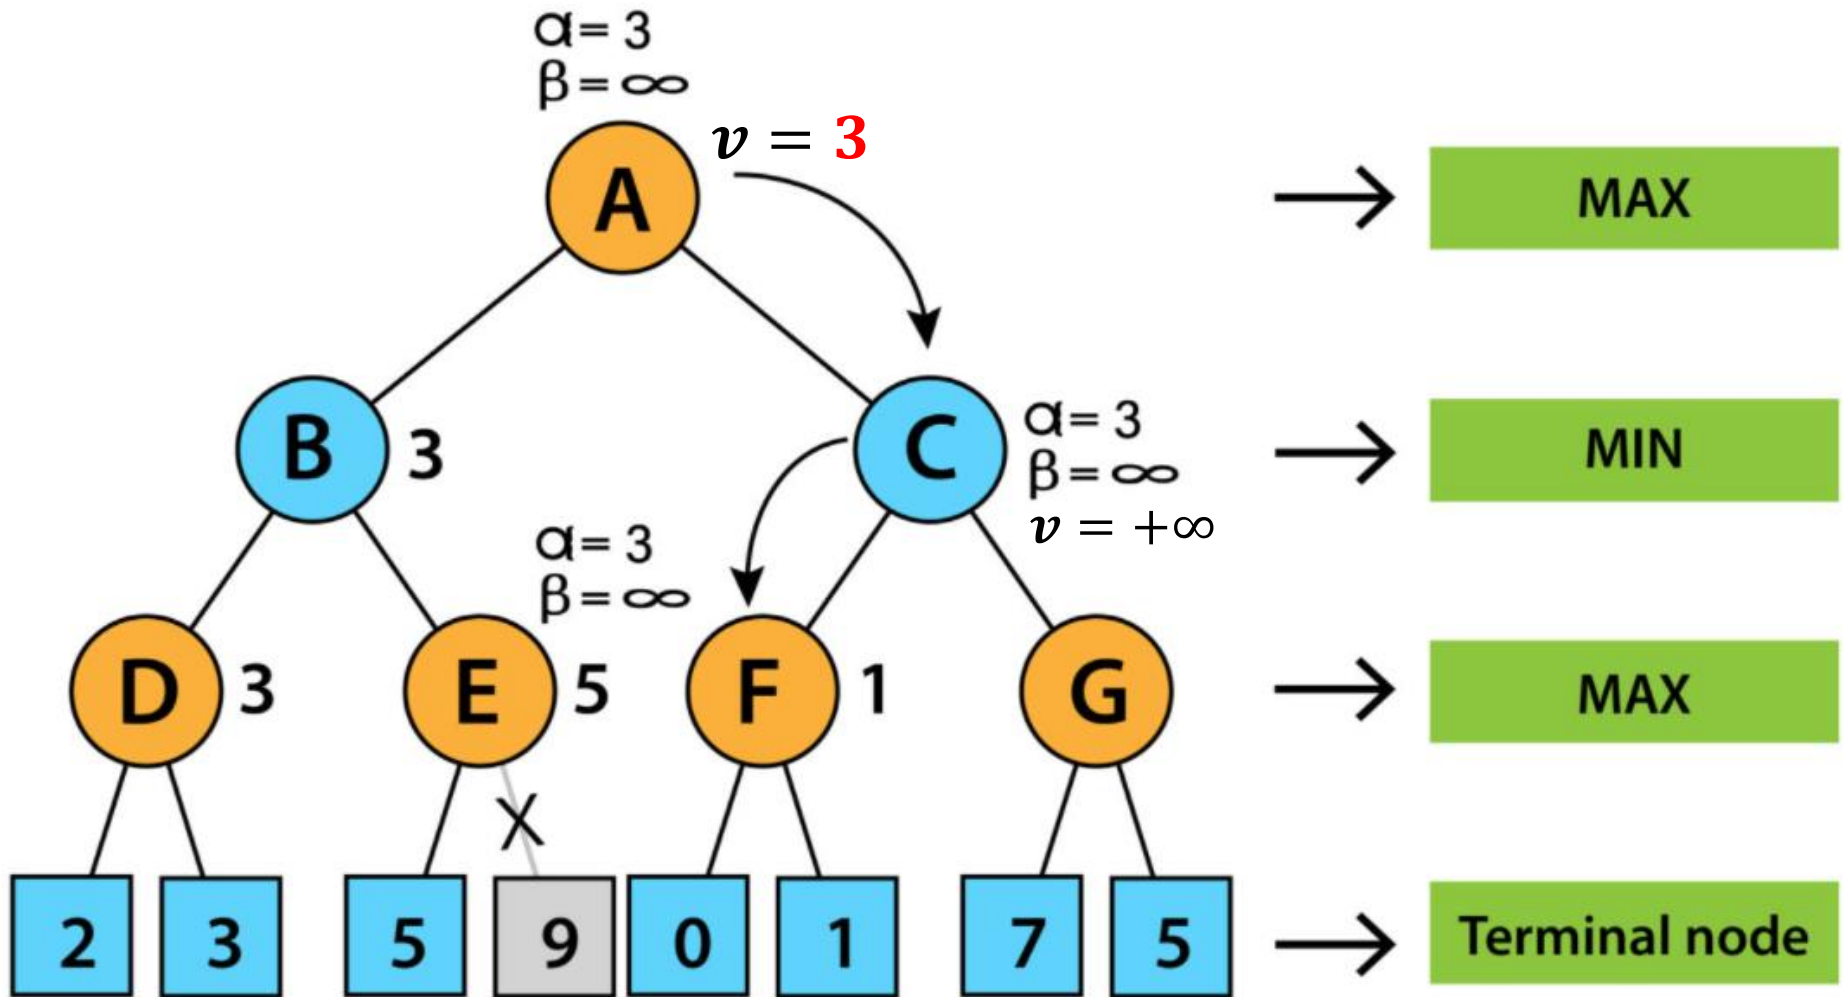
Alpha-Beta Pruning Example



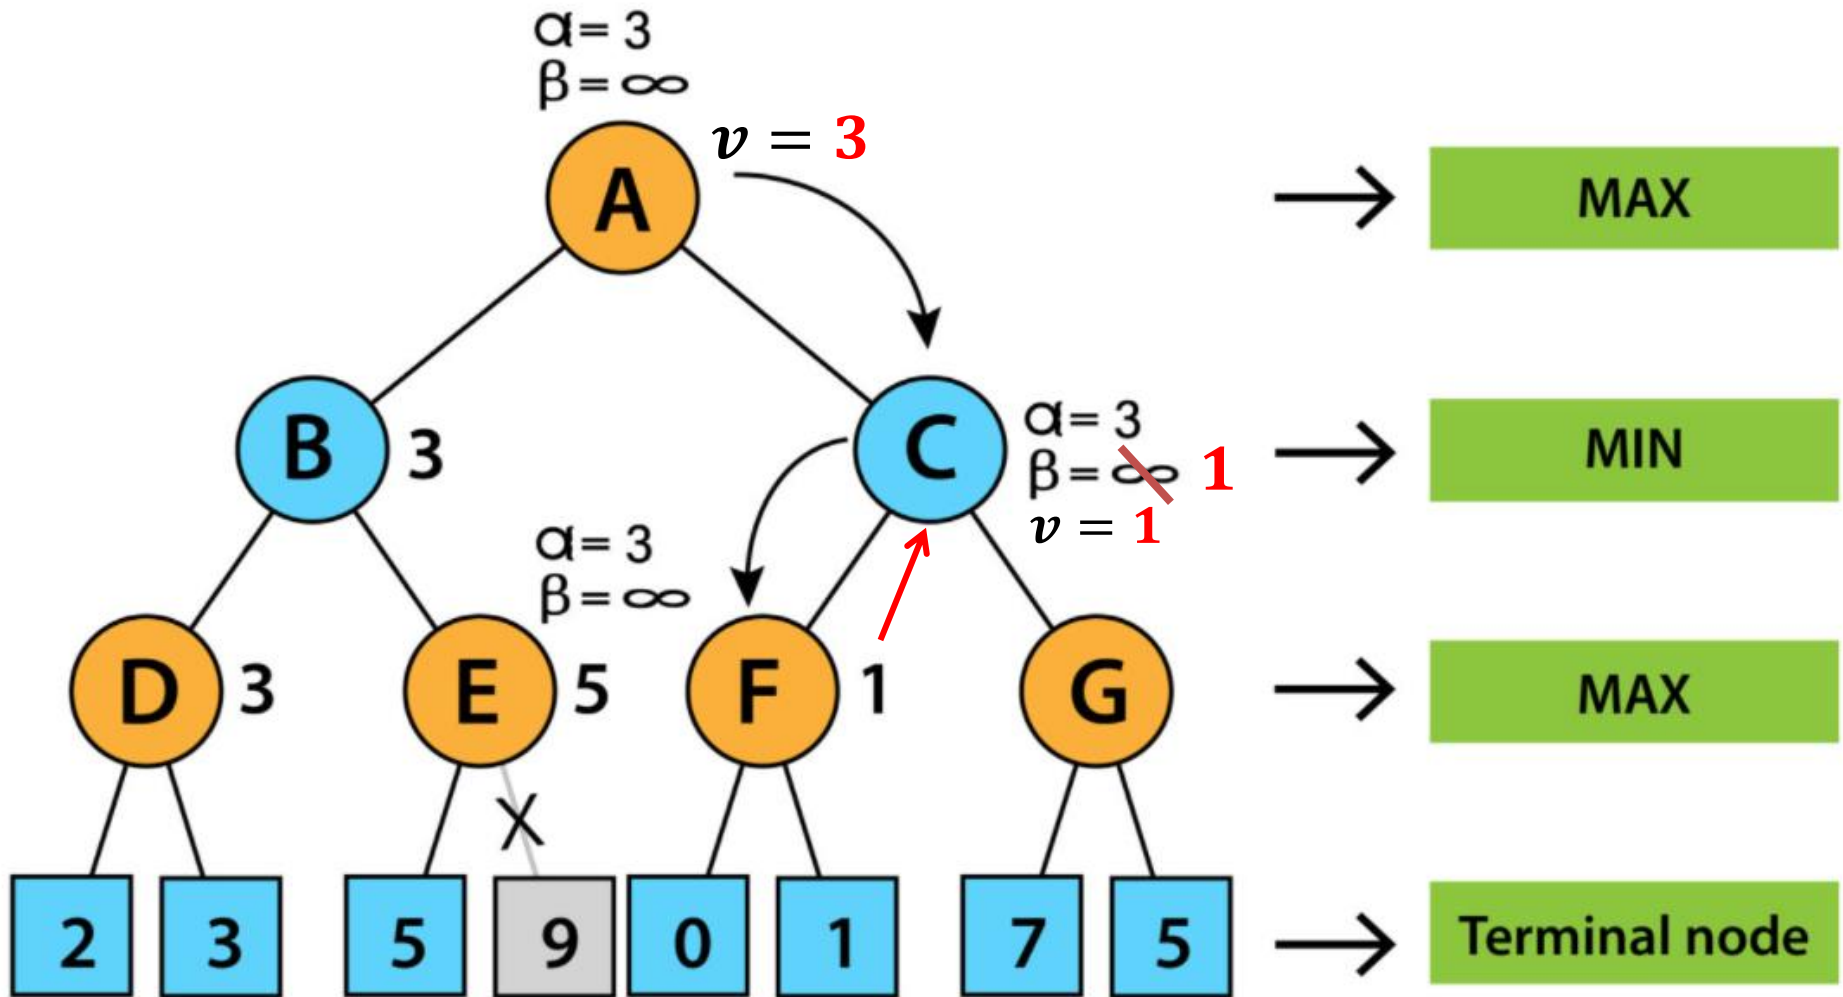
Alpha-Beta Pruning Example



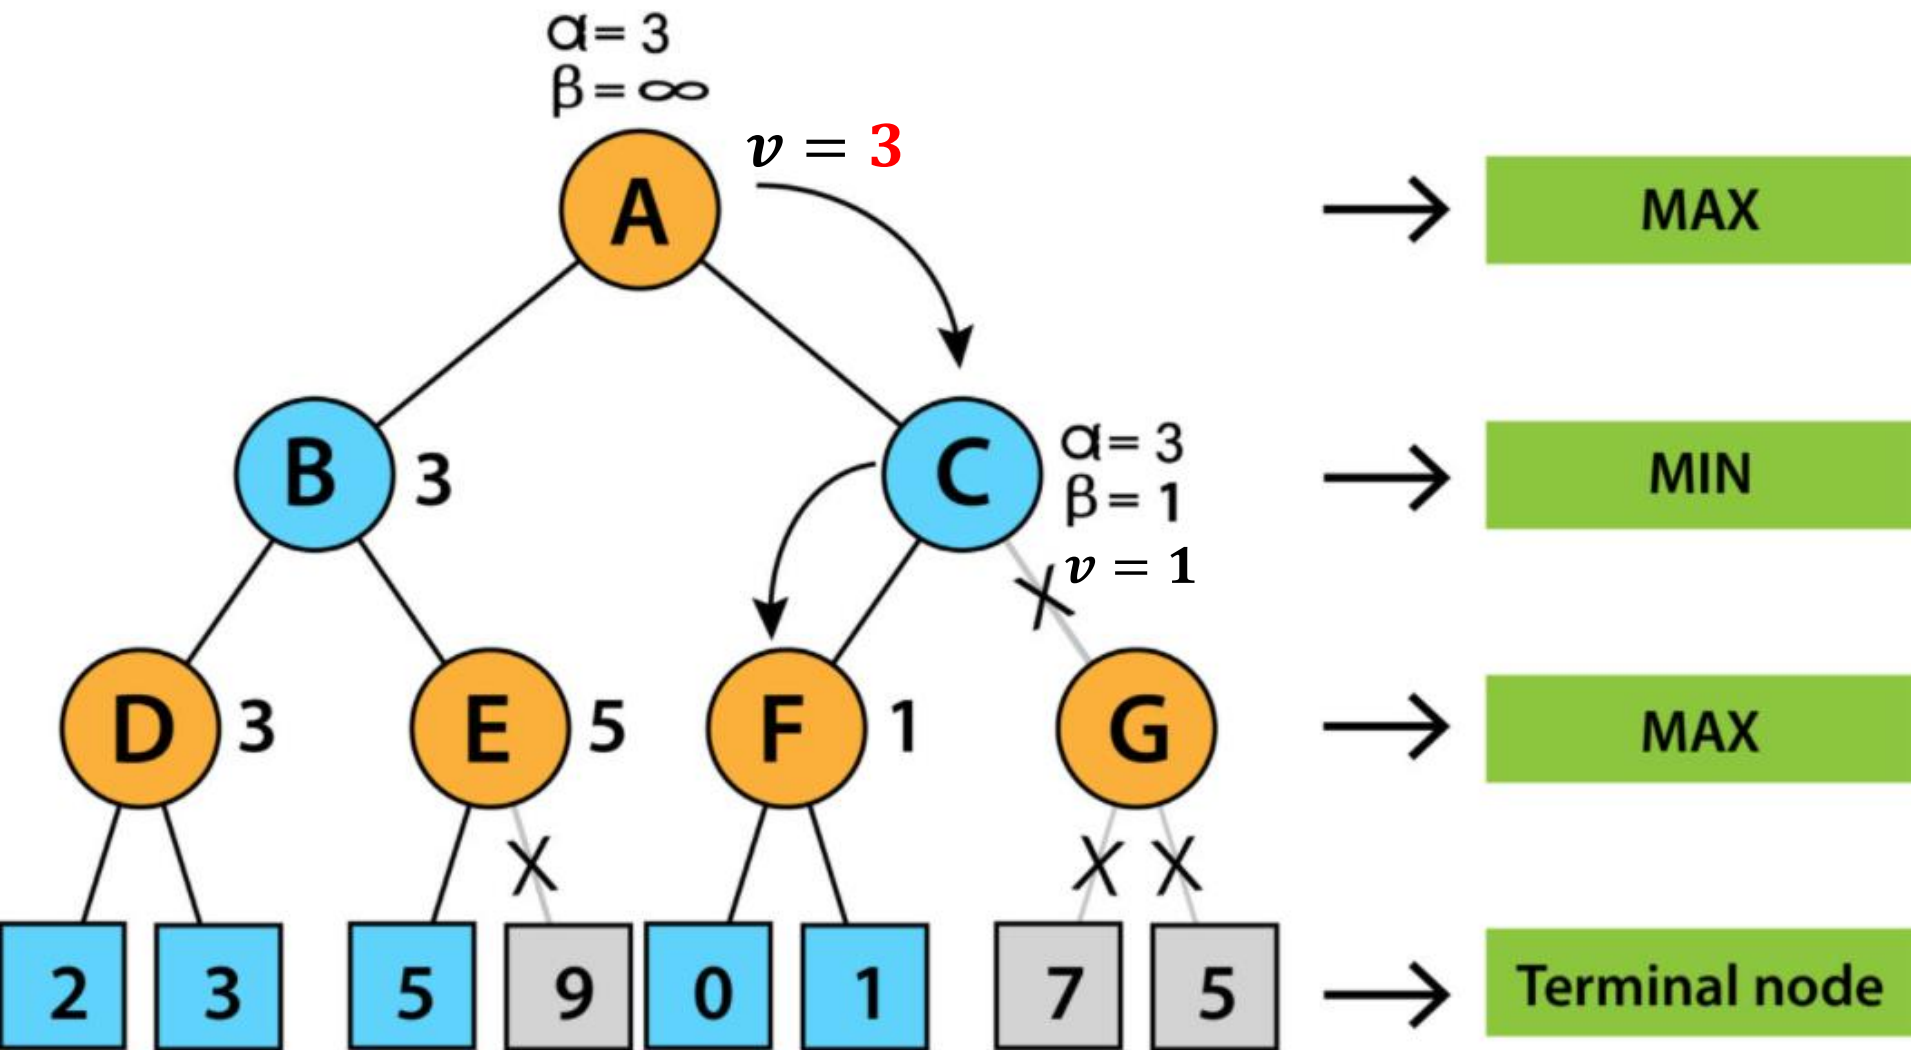
Alpha-Beta Pruning Example



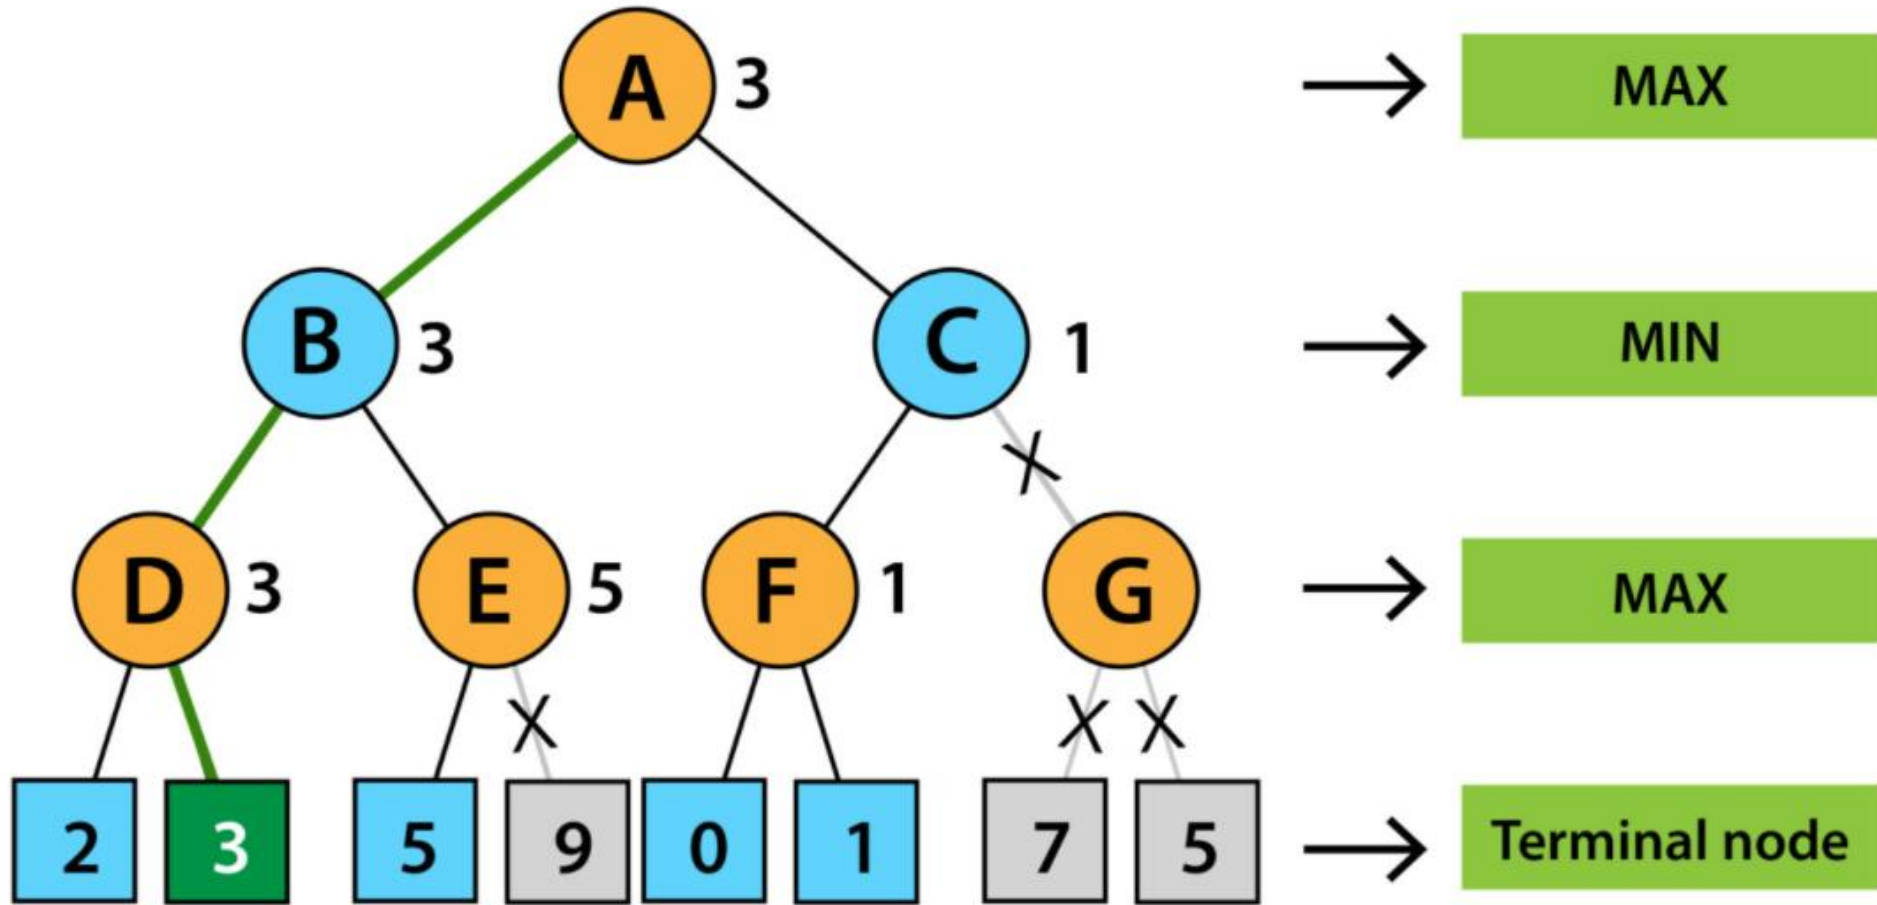
Alpha-Beta Pruning Example



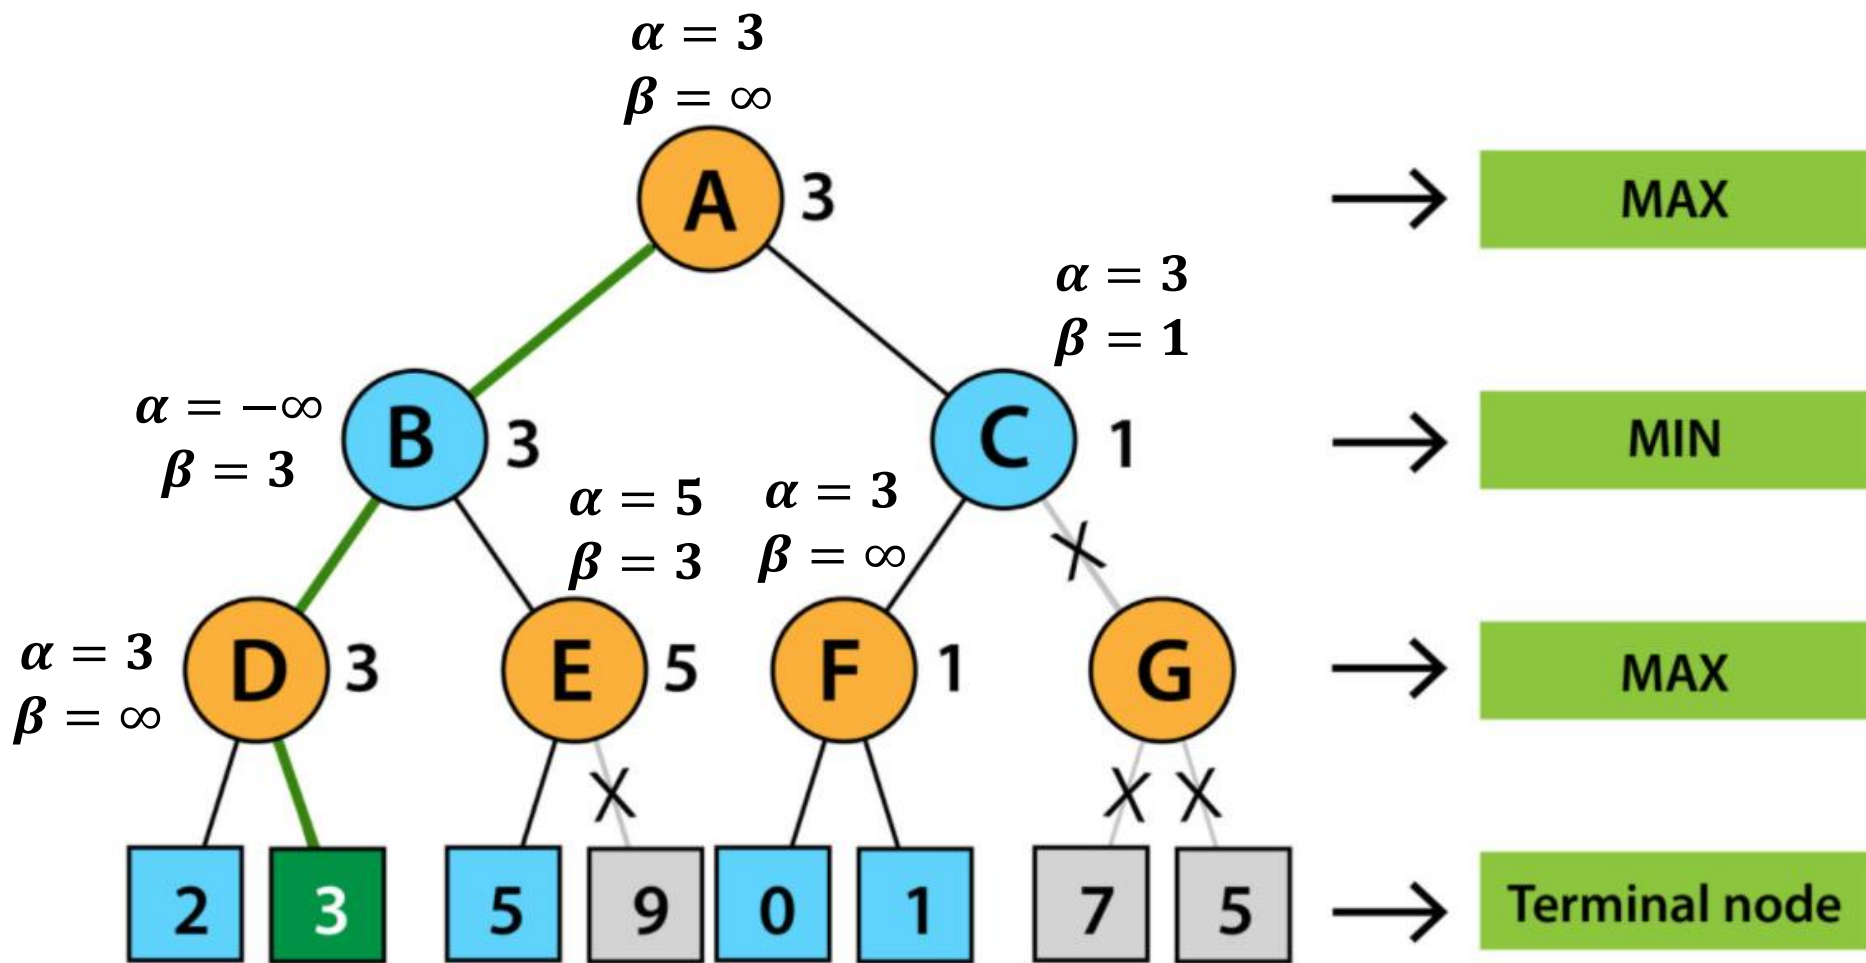
Alpha-Beta Pruning Example




Alpha-Beta Pruning Example




Alpha-Beta Pruning Example




Exercise



Exercise: Solution



Exercise



Properties of Alpha-Beta Algorithm

- Pruning has no effect on the minimax values.
- Entire subtrees can be pruned, not just leaves.
- Good move ordering improves effectiveness of pruning
- Completeness and optimality are preserved from Minimax.
- Time complexity:
 - The effectiveness depends on the order in which the states are examined.
 - If states could be examined in **perfect order**, then alpha-beta search examines only $O(b^{m/2})$ nodes to pick the best move, vs. $O(b^m)$ for minimax.
 - If states could be examined in **random order** rather than perfect order, the total number of nodes examined will be roughly $O(b^{3m/4})$ for moderate b .
- Space complexity: $O(m)$, as for Minimax.

Game Tree Size



E.g., Chess:

- $b \approx 35$ (approximate average branching factor)
- $d \approx 100$ (depth of a game tree for typical games)
- $b^d \approx 35^{100} \approx 10^{154}$
- For alpha-beta search and perfect ordering, we still get $b^{d/2} \approx 35^{50} \approx 10^{77}$

Finding the exact solution is completely **infeasible!**

Is this practical?

- Minimax and alpha-beta pruning still have exponential complexity.
- May be impractical within a reasonable amount of time.
- Solution: cut the search earlier
 - replace the **Utility** function with a **heuristic** evaluation function that estimates the state utility.
 - replace the **Terminal-test** by a **cutoff test** that decides when to stop expanding a state.

Cutting off search

- Introduces a fixed-depth limit d
 - Selected so that the amount of time will not exceed what the rules of the game allow.
- When cutoff occurs, the evaluation is performed.
- Heuristic MiniMax for state s and maximum depth d

```
H-MiniMax( $s$ ,  $d$ ) =  
  
if Cutoff-Test( $s$ ,  $d$ ) then Eval( $s$ )  
if Player( $s$ ) = MAX then  
    max of H-MiniMax(Result( $s$ ,  $a$ ),  $d+1$ ) for  $a$  in Actions( $s$ )  
if Player( $s$ ) = MIN then  
    min of H-MiniMax(Result( $s$ ,  $a$ ),  $d+1$ ) for  $a$  in Actions( $s$ )
```

Evaluation Function

- Idea: produce an estimate of the expected utility of the game from a given position.
- Performance depends on quality of Eval.
- Requirements:
 - must be consistent with the utility function
 - states that are wins must evaluate better than draws, which in turn must be better than losses.
 - tradeoff between accuracy and time cost
 - without time limits, minimax could be used
 - should reflect the actual **chances of winning** for non-terminal states
- Frequently weighted linear functions are used
 - $E = w_1 f_1 + w_2 f_2 + \dots + w_n f_n$
 - combination of features (f_1, \dots, f_n), weighted by their relevance (w_1, \dots, w_n)
 - more important features get more weight

The horizon effect

- Evaluations functions are always imperfect.
- If not looked deep enough, bad moves may appear as good moves (as estimated by the evaluation function) because their consequences are hidden beyond the search horizon.
 - and vice-versa!
- Often, the deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters.

Summary

- **Game Tree**
- **Alpha-Beta Pruning**
- **Different Games**

Midterm Exam

- 0am 02/29 – 11:59pm 03/01 under “Exams” of Modules on Canvas
- Time duration: 3 hours
- Open-book and open-note
- Total points: 100
- Five Problems
- 1) Search problem: BFS, DFS, UCS, Greedy, A*
- 2) Search Problem
- 3) Local search- simulated annealing
- 4) Constraint satisfaction problem
- 5) Games: minimax

What I want you to do

- Review Chapter 5
- Work on your assignment