

Lecture 21

Learning II

Lusi Li

**Department of Computer Science
ODU**

Reading for This Class:
Not in the Book

Review

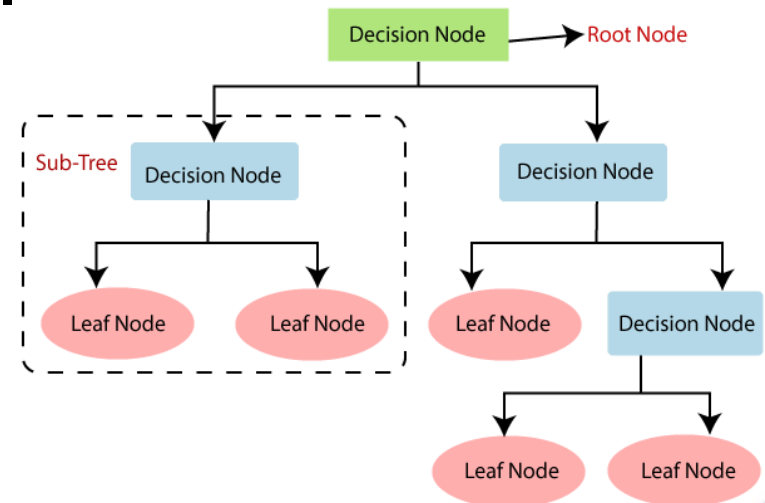
- **Last Class**
 - Learning
 - Decision Tree
- **This Class**
 - Nearest Neighbors Classification
 - Unsupervised Learning
 - K-means
- **Next Class**
 - Final Review

Basic Algorithm of Decision Tree

node = root of decision tree

Main loop:

1. $A \leftarrow$ the “**best**” decision attribute for the next node.
2. Assign A as decision attribute for node.
3. For each value of A , create a new child (sub-tree) of the node.
4. Sort training examples to leaf nodes.
5. If training examples are perfectly classified, stop.
Else, recurse over new leaf nodes.



Disadvantages of Decision-Tree Learning

- **Large-Scale Information**
 - A big decision tree
 - Not efficient
- **Contradictory information**
 - Fail to build a decision tree
 - Decision tree is not robust to contradictory or erroneous information
 - Hard to handle noisy information
- **Missing Information**
 - Not all the attributed values are known in some given examples
 - Hard to classify
- **Adaptability**
 - Learning Decision Tree may not be useful in a changing environment
- **Real Time Response**
 - If the decision tree is large, response time may be long

Outline

- Nearest Neighbor Classification
- Intro to unsupervised learning
- K-means algorithm
- Optimization objective
- Initialization and the number of clusters

Supervised learning

- **Input:** training set (input-output pairs):
 - $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$
 - where each pair was generated by an unknown function f

$$f(x^{(i)}) = y^{(i)}, 1 \leq i \leq m$$

- **Goal:** find a hypothesis function h that approximates the true function f


$$h(x^{(i)}) \approx y^{(i)}, 1 \leq i \leq m$$

What we know so far

- Decision Trees: how to induce a decision tree from training data
- Other methods?


KNN

- K-Nearest Neighbor (KNN) Classifier
 - Organize and store all training examples
 - Classify new examples based on “most similar” training examples
 - Compute distance to other training examples
 - Identify k nearest neighbors
 - Use class labels of KNN to determine the class label by taking majority vote




$k = 1$

Nearest point is **red**, so
 x_{test} classified as **red**



$k = 3$

Nearest points are {**red**,
blue, **blue**} so x_{test}
classified as **blue**



$k = 4$

Nearest points are {**red**,
red, **blue**, **blue**} so
classification of x_{test} is
not properly defined

Supervised learning

- Two ways to think about learning

Eager learning (e.g., decision trees)

- Learn/Train
 - Induce an **abstract model** from data
- Test/Predict/Classify
 - Apply learned model to new data

Lazy learning (e.g., nearest neighbors)

- Learn
 - **Just store data** in memory
- Test/Predict/Classify
 - Compare new data to stored data
- **Properties**
 - Retains all information seen in training
 - Complex hypothesis space
 - Classification can be very slow

KNN algorithm


Components of a k-NN Classifier

- **Distance metric**
 - How do we measure distance between instances?
 - Determines the layout of the example space
- **The k hyper-parameter**
 - How large a neighborhood should we consider?
 - Determines the complexity of the hypothesis space

KNN algorithm


- **Distance metric**
 - Any distance function can select nearest neighbors
 - Different distances yield different neighborhoods

Distance Metrics in Vector Search




Cosine Distance

$$1 - \frac{A \cdot B}{||A|| \ ||B||}$$




Squared Euclidean
(L2 Squared)

$$\sum_{i=1}^n (x_i - y_i)^2$$



Dot Product

$$A \cdot B = \sum_{i=1}^n A_i B_i$$



Manhattan (L1)


$$\sum_{i=1}^n |x_i - y_i|$$

KNN algorithm

- **The k hyper-parameter**

- If k is too small, sensitive to noise points
- If k is too large, neighborhood may include points

What class does the new data point belong to?



How would you set k in practice?

- **Weighted voting**
 - Default:
all neighbors have equal weight
 - Extension:
weight neighbors by (inverse)
distance

Unsupervised learning

- **Input:** training set: $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots, x^{(m)}\}$
- **Clustering goal:** automatically partition examples into groups of similar examples
- Why? It is useful for
 - Automatically organizing data
 - Understanding hidden structure in data
 - Preprocessing for further analysis

K-means algorithm

- Input:
 - K (number of clusters)
 - Training set $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots, x^{(m)}\}$, where $x^{(i)} \in \mathbb{R}^n$
- Randomly initialize K cluster centroids $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^n$

Repeat{

for $i = 1$ to m

$c^{(i)} := \text{index (from 1 to } K) \text{ of cluster centroid closest to } x^{(i)}$

Cluster assignment step

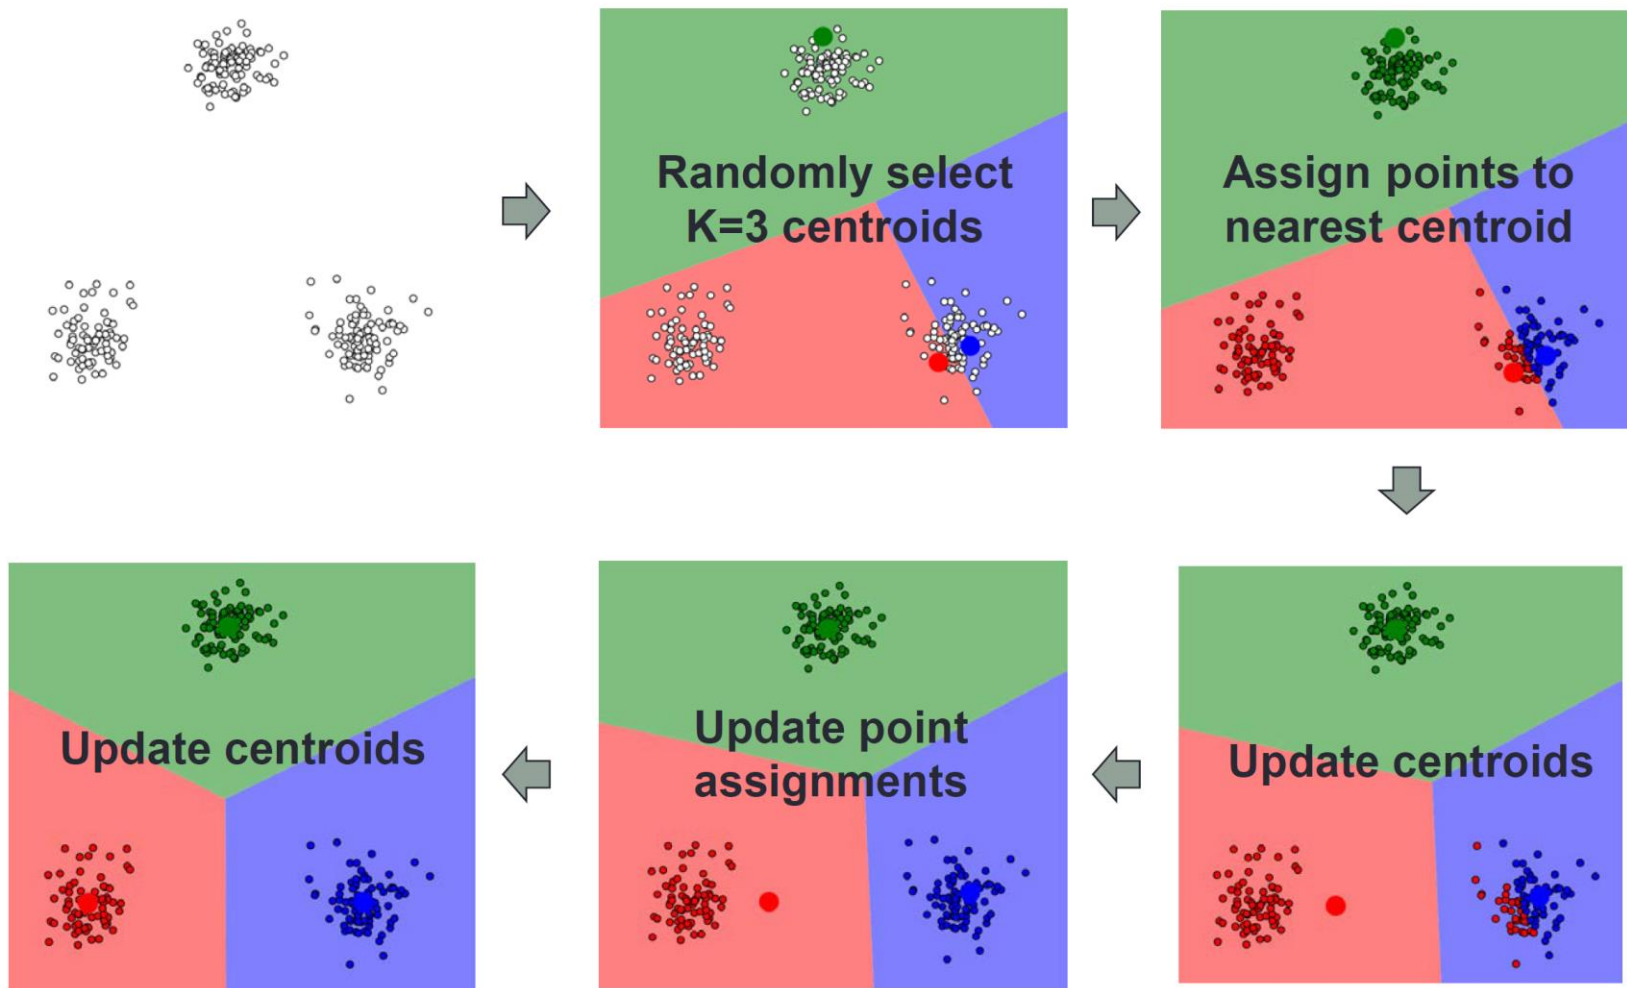
for $k = 1$ to K

$\mu_k := \text{average (mean) of points assigned to cluster } k$

Centroid update step

}

K-means algorithm



K-Means example

From <https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

K-means optimization objective

- $c^{(i)}$ = Index of cluster (1, 2, ... K) to which example $x^{(i)}$ is currently assigned
- μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$)
- $\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned
- Optimization objective:

Example:

For $x^{(i)} \in \mathbb{R}^n$

$$c^{(i)} = 2$$

$$\mu_{c^{(i)}} = \mu_2$$

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{\substack{c^{(1)}, \dots, c^{(m)} \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

K-means algorithm

Randomly initialize K cluster centroids $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^n$

Repeat{

for $i = 1$ to m

Cluster assignment step

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$c^{(i)} :=$ index (from 1 to K) of cluster centroid
closest to $x^{(i)}$

Centroid update step

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

for $k = 1$ to K

$\mu_k :=$ average (mean) of points assigned to cluster k

}

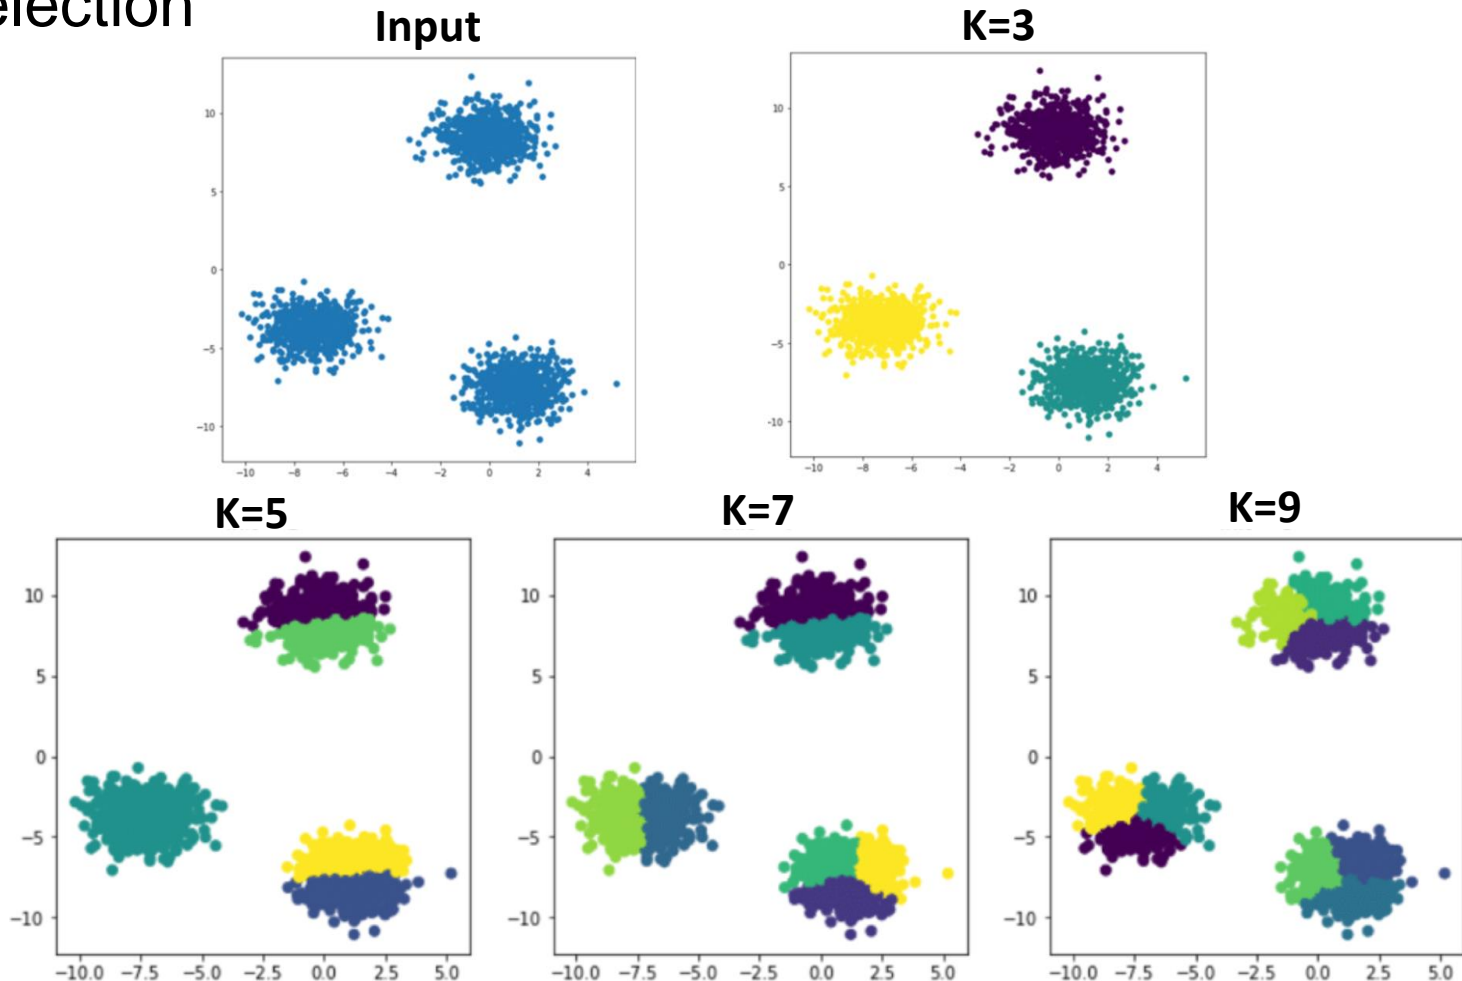
K-means algorithm

Components of K-means

- **Distance metric**
 - How do we measure distance between instances?
 - Determines the cluster assignment
- **The K hyper-parameter**
 - Needs to be set in advance (as prior knowledge)
- **Cluster-centroid initialization**
 - Different initializations yield different results


How to choose K?

- Try multiple **K** and use performance from downstream task for selection



Impact of initialization

- Randomly pick K training examples
- Set $\mu_1, \mu_2, \dots, \mu_K$ equal to those K examples



1) Multiple random initialization

For $i = 1$ to 100 {

Randomly initialize K-means

Run K-means. Get $c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K$

Compute the cost function (distortion)

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

}

Pick clustering that gave the lowest cost

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

2) Furthest point heuristic

Choose μ_1 arbitrarily (or at random)

For $j = 2$ to K

Pick μ_j among data points $x^{(1)}, x^{(2)}, \dots, x^{(m)}$ that is farthest from previously chosen $\mu_1, \mu_2, \dots, \mu_{j-1}$

Slide credit: Maria-Florina Balcan

Things to remember

- Nearest Neighbor Classification
- Intro to unsupervised learning
- K-means algorithm
- Optimization objective
- Initialization and the number of clusters

What I want you to do

- **Work on your assignments 3 and 4**